

# Going beyond linear regression

GENERALIZED LINEAR MODELS IN PYTHON

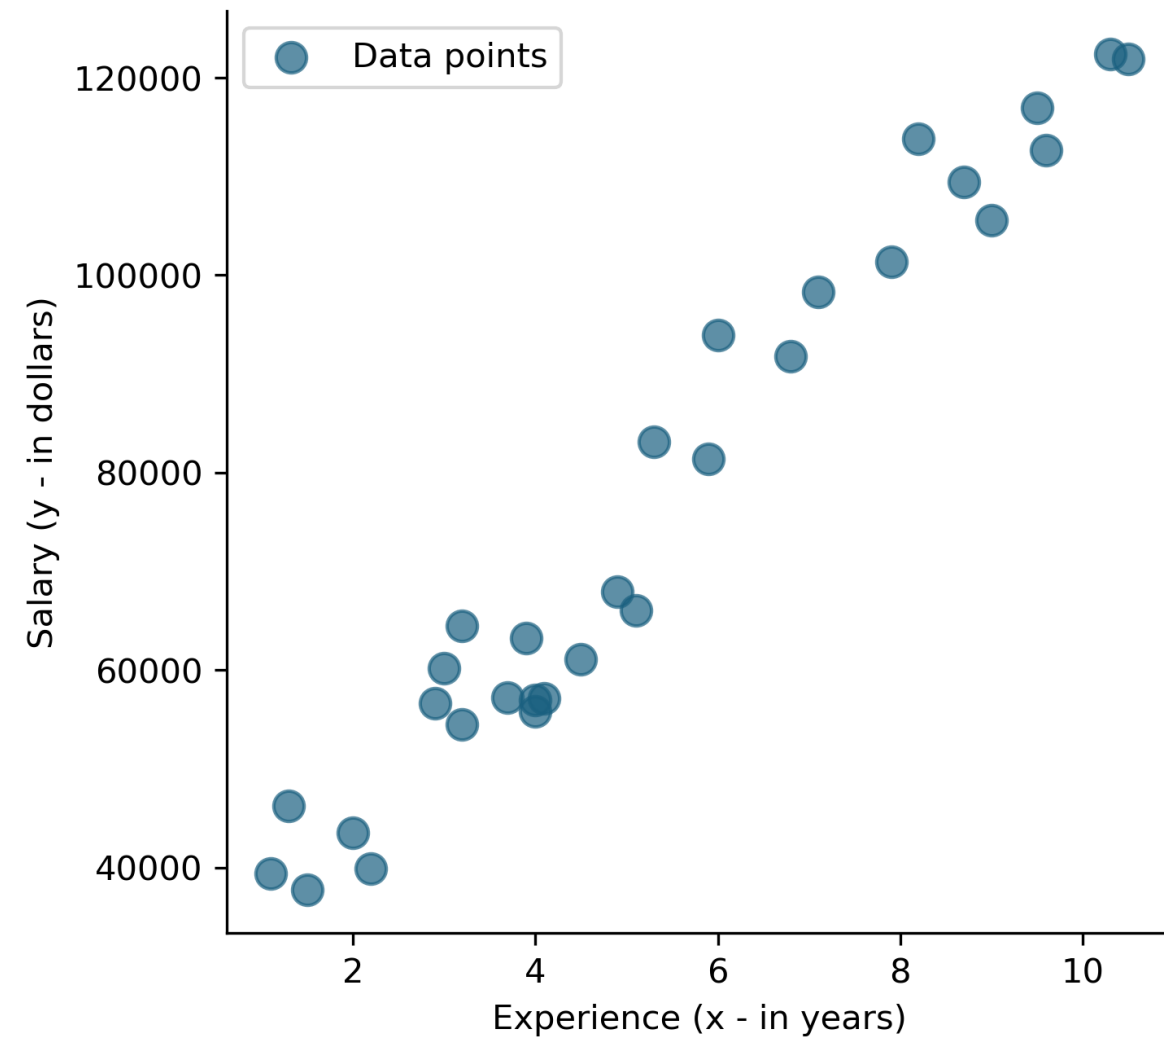


**Ita Cirovic Donev**  
Data Science Consultant

# Course objectives

- Learn building blocks of GLMs
  - Train GLMs
  - Interpret model results
  - Assess model performance
  - Compute predictions
- Chapter 1: How are GLMs an extension of linear models
  - Chapter 2: Binomial (logistic) regression
  - Chapter 3: Poisson regression
  - Chapter 4: Multivariate logistic regression

# Review of linear models

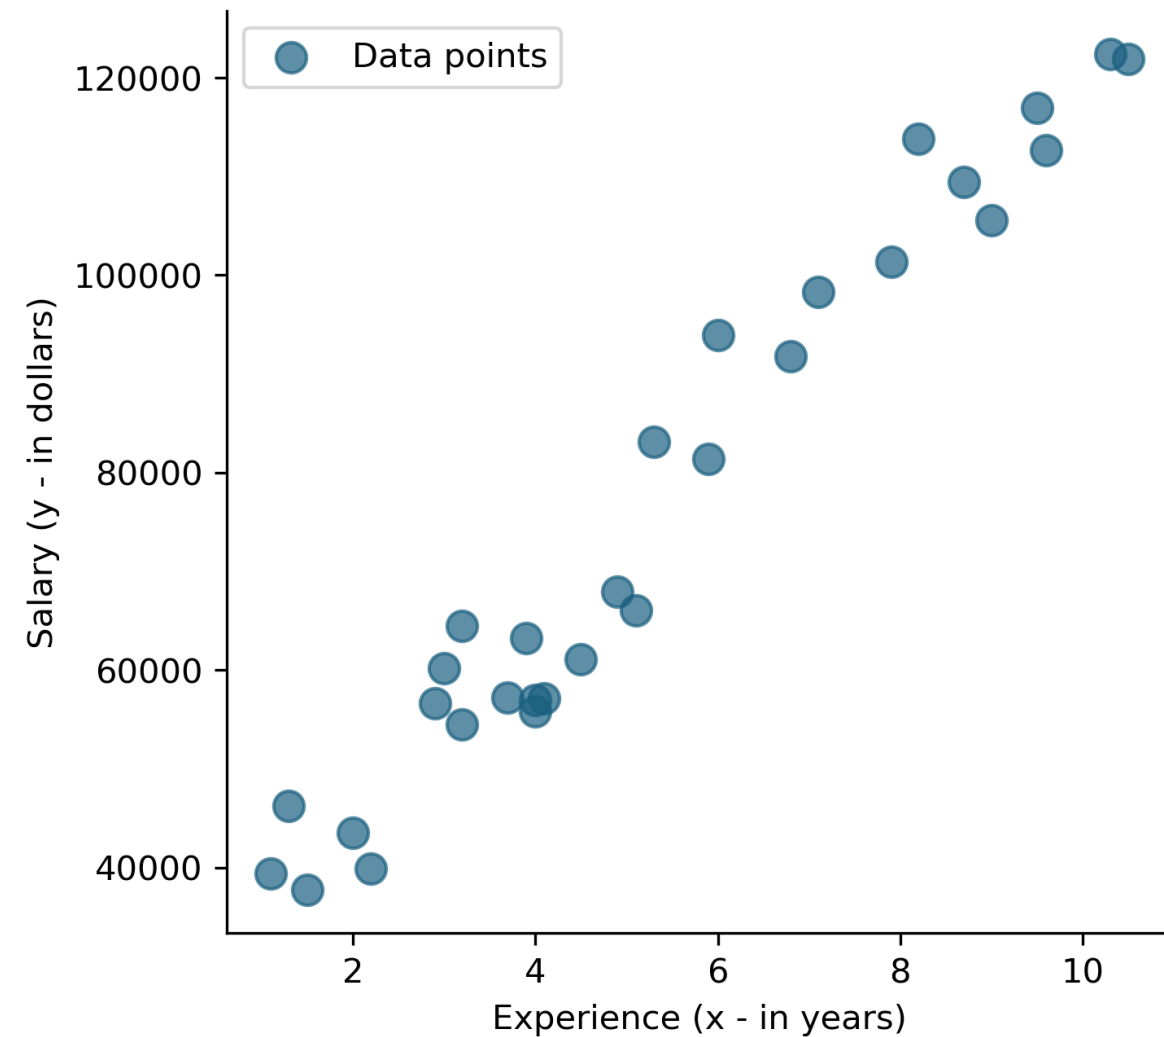


salary  $\sim$  experience

salary  $= \beta_0 + \beta_1 \times \text{experience} + \epsilon$

$y = \beta_0 + \beta_1 x_1 + \epsilon$

# Review of linear models



`salary`  $\sim$  `experience`


$$\text{salary} = \beta_0 + \beta_1 \times \text{experience} + \epsilon$$

$$y = \beta_0 + \beta_1 x_1 + \epsilon$$

where:

`y` - response variable (output)

# Review of linear models



$\text{salary} \sim \text{experience}$

$\text{salary} = \beta_0 + \beta_1 \times \text{experience} + \epsilon$


$y = \beta_0 + \beta_1 x_1 + \epsilon$

where:

$y$  - response variable (output)

$x$  - explanatory variable (input)

# Review of linear models



salary  $\sim$  experience

$$\text{salary} = \beta_0 + \beta_1 \times \text{experience} + \epsilon$$

$$y = \beta_0 + \beta_1 x_1 + \epsilon$$

where:

$y$  - response variable (output)

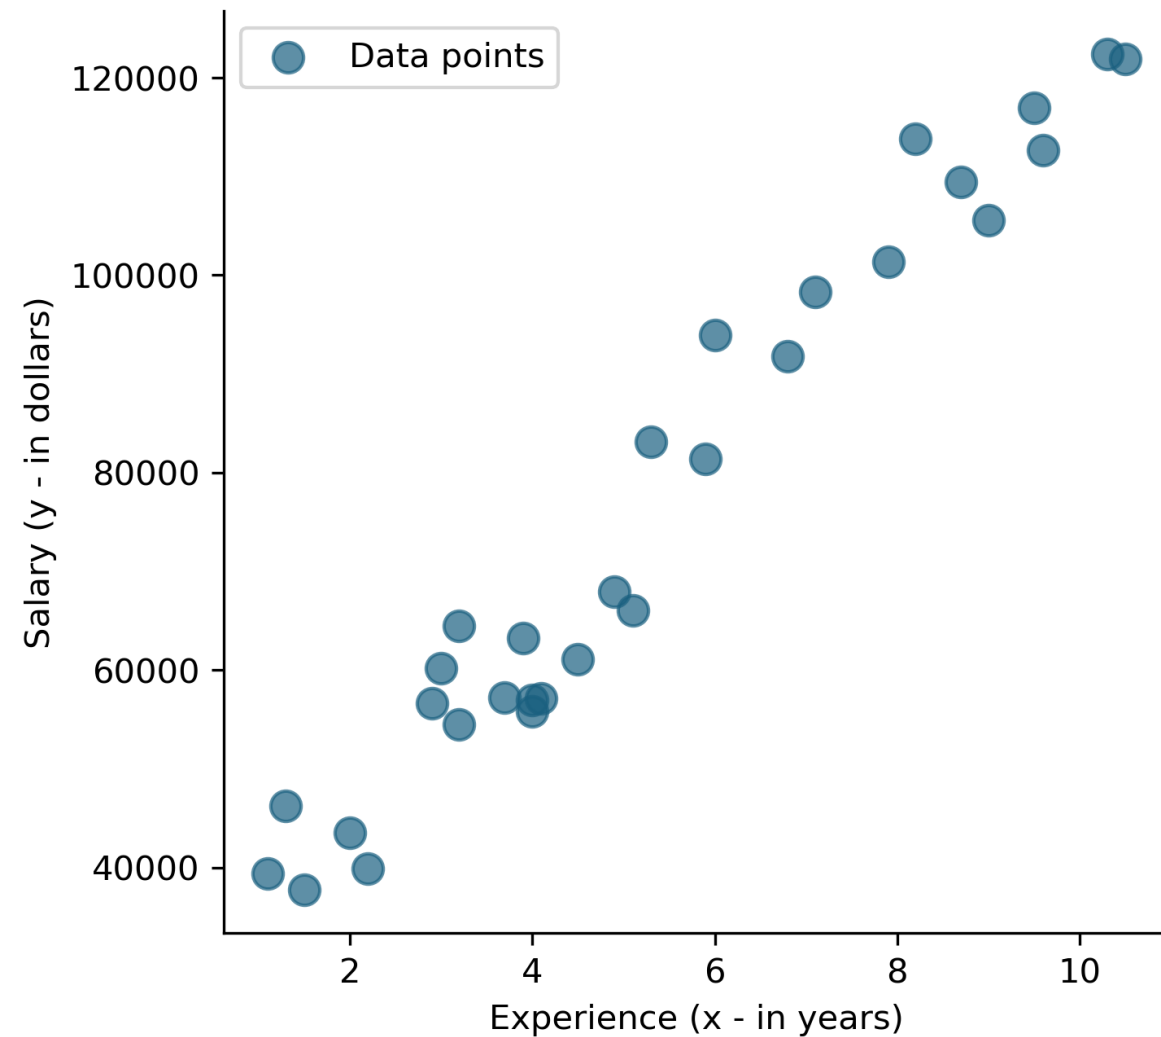
$x$  - explanatory variable (input)

$\beta$  - model parameters

$\beta_0$  - intercept

$\beta_1$  - slope

# Review of linear models



salary  $\sim$  experience

$$\text{salary} = \beta_0 + \beta_1 \times \text{experience} + \epsilon$$

$$y = \beta_0 + \beta_1 x_1 + \epsilon$$

where:

$y$  - response variable (output)

$x$  - explanatory variable (input)

$\beta$  - model parameters

$\beta_0$  - intercept

$\beta_1$  - slope

$\epsilon$  - random error

## LINEAR MODEL - `ols()`

```
from statsmodels.formula.api import ols
```

```
model = ols(formula = 'y ~ X',  
             data = my_data).fit()
```

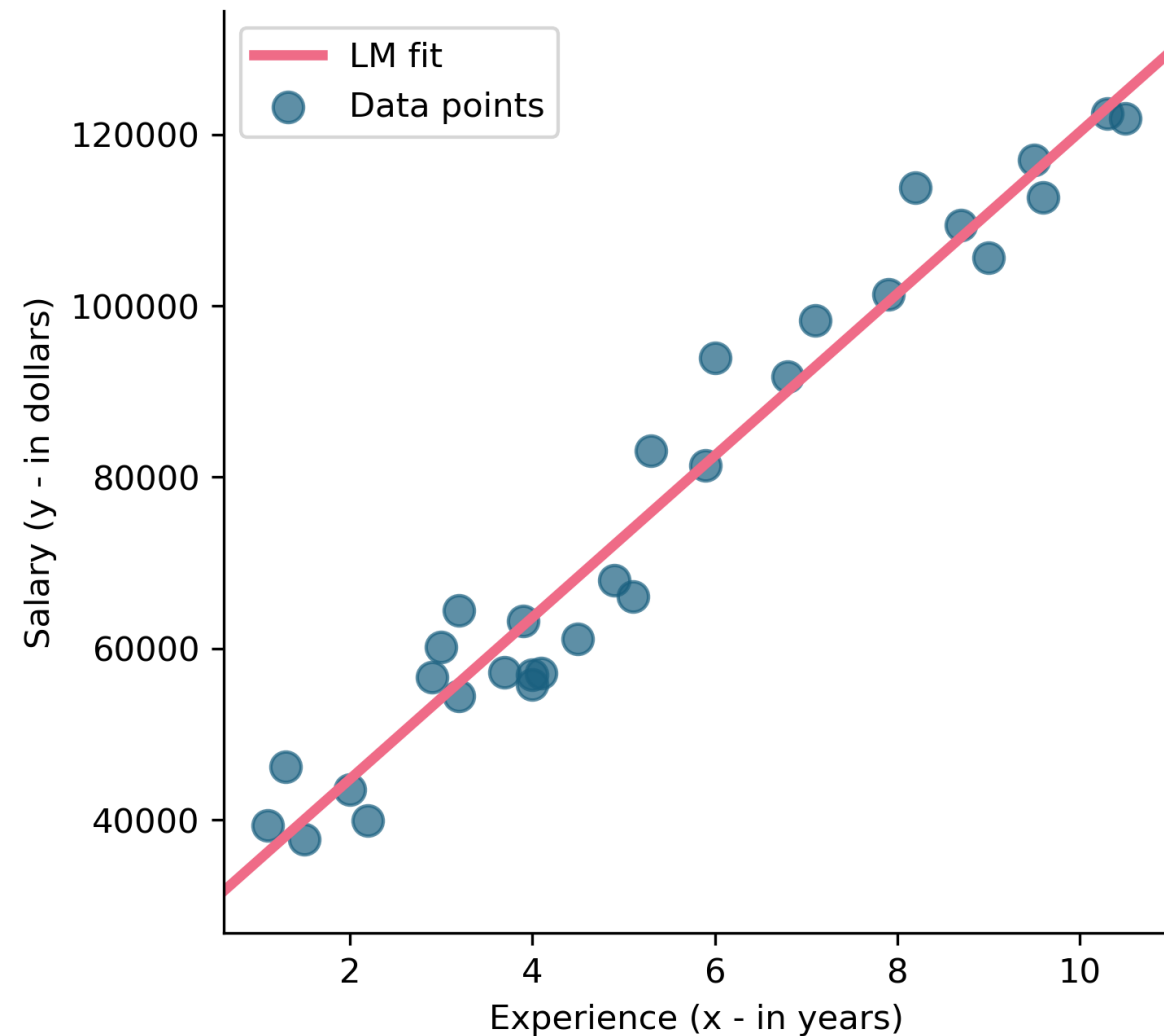
## GENERALIZED LINEAR MODEL - `glm()`

```
import statsmodels.api as sm  
from statsmodels.formula.api import glm
```

```
model = glm(formula = 'y ~ X',  
            data = my_data,  
            family = sm.families.Gaussian).fit
```



# Assumptions of linear models



## Regression function

$$E[y] = \mu = \beta_0 + \beta_1 x_1$$

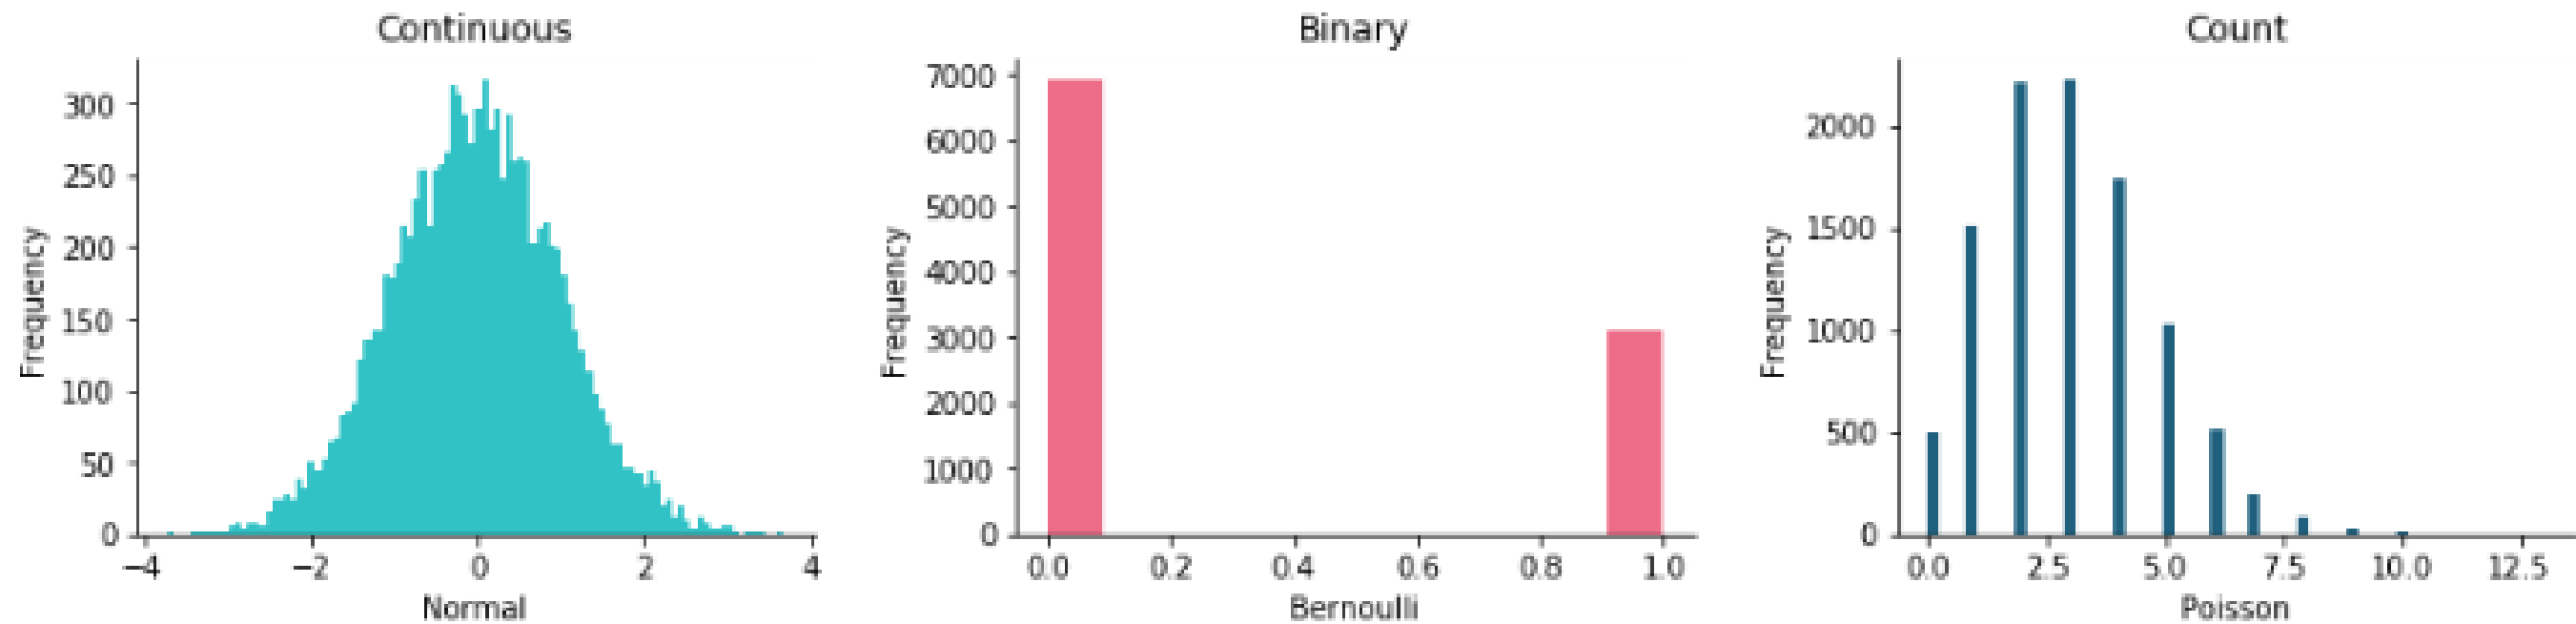
## Assumptions

- Linear in parameters
- Errors are independent and normally distributed
- Constant variance

$$\text{salary} = 25790 + 9449 \times \text{experience}$$

# What if ... ?

- The response is binary or count → **NOT continuous**



- The variance of  $y$  is not constant → **depends on the mean**

# Dataset - nesting of horseshoe crabs

Variable Name	Description
sat	Number of satellites residing in the nest
y	There is at least one satellite residing in the nest; 0/1
weight	Weight of the female crab in kg
width	Width of the female crab in cm
color	1 - light medium, 2 - medium, 3 - dark medium, 4 - dark
spine	1 - both good, 2 - one worn or broken, 3 - both worn or broken

<sup>1</sup> A. Agresti, An Introduction to Categorical Data Analysis, 2007.

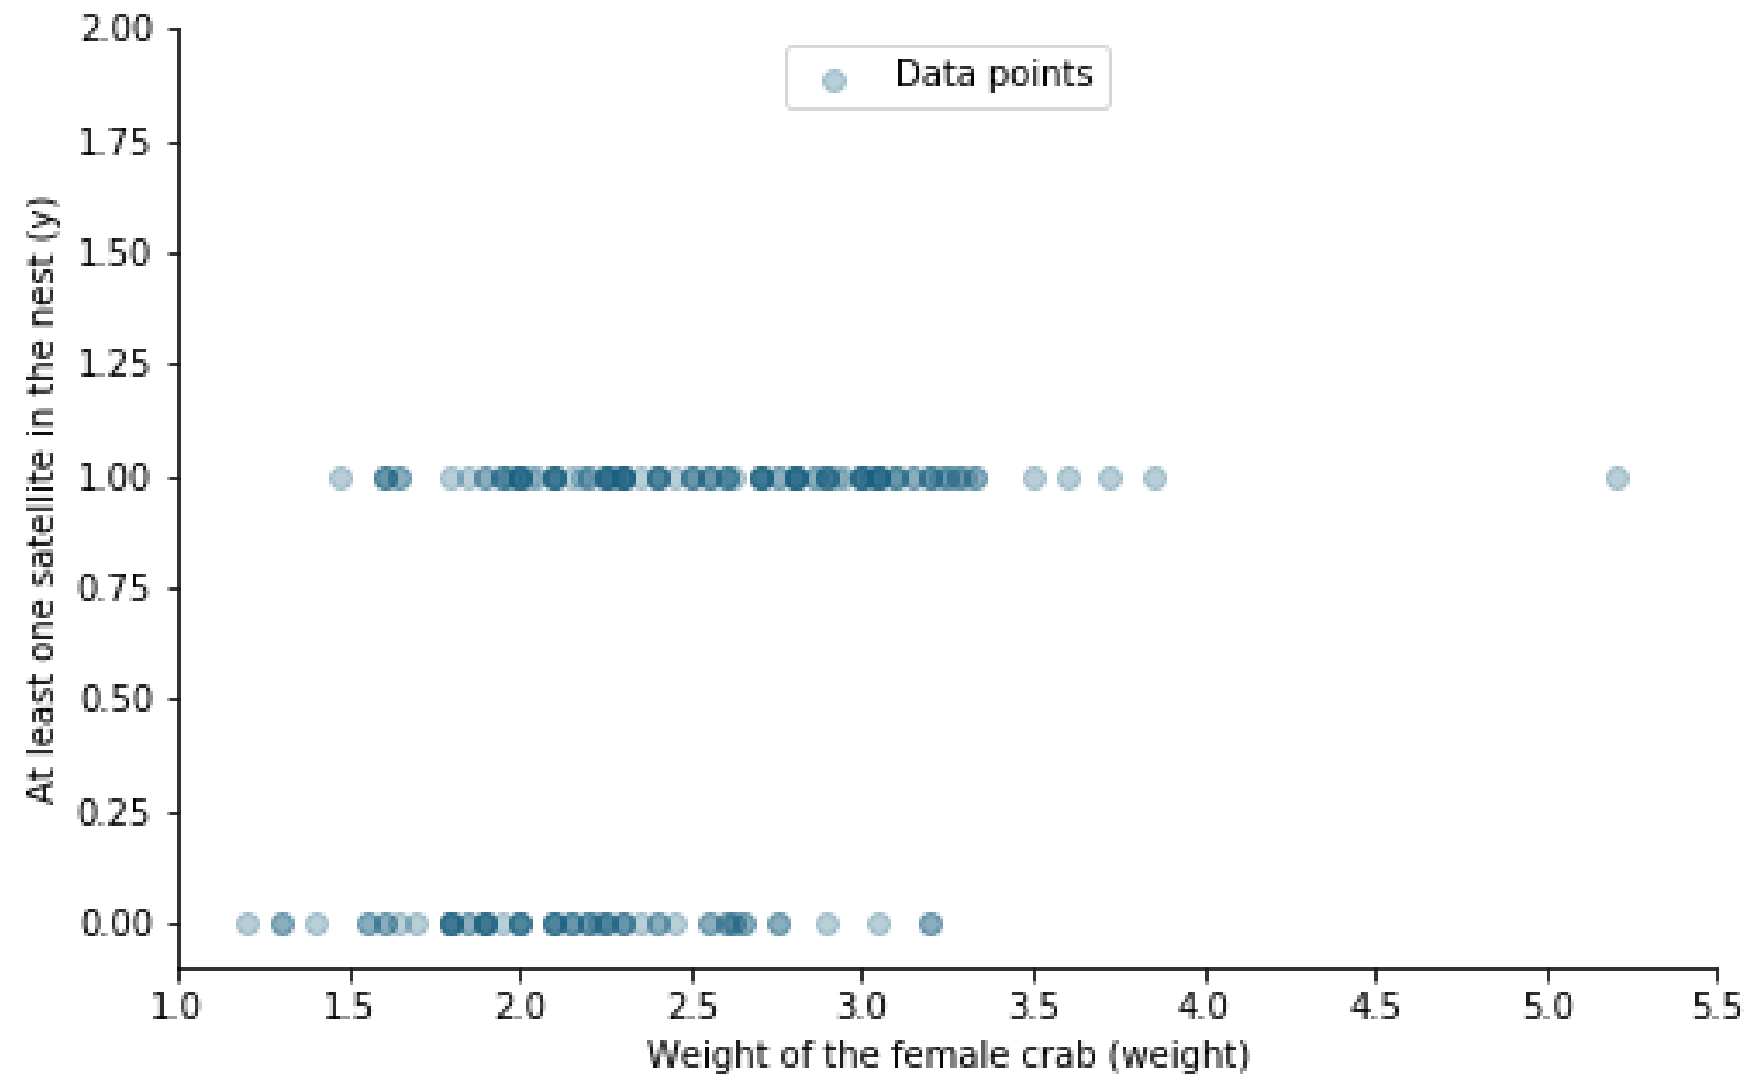
# Linear model and binary response

satellite crab  $\sim$  female crab weight

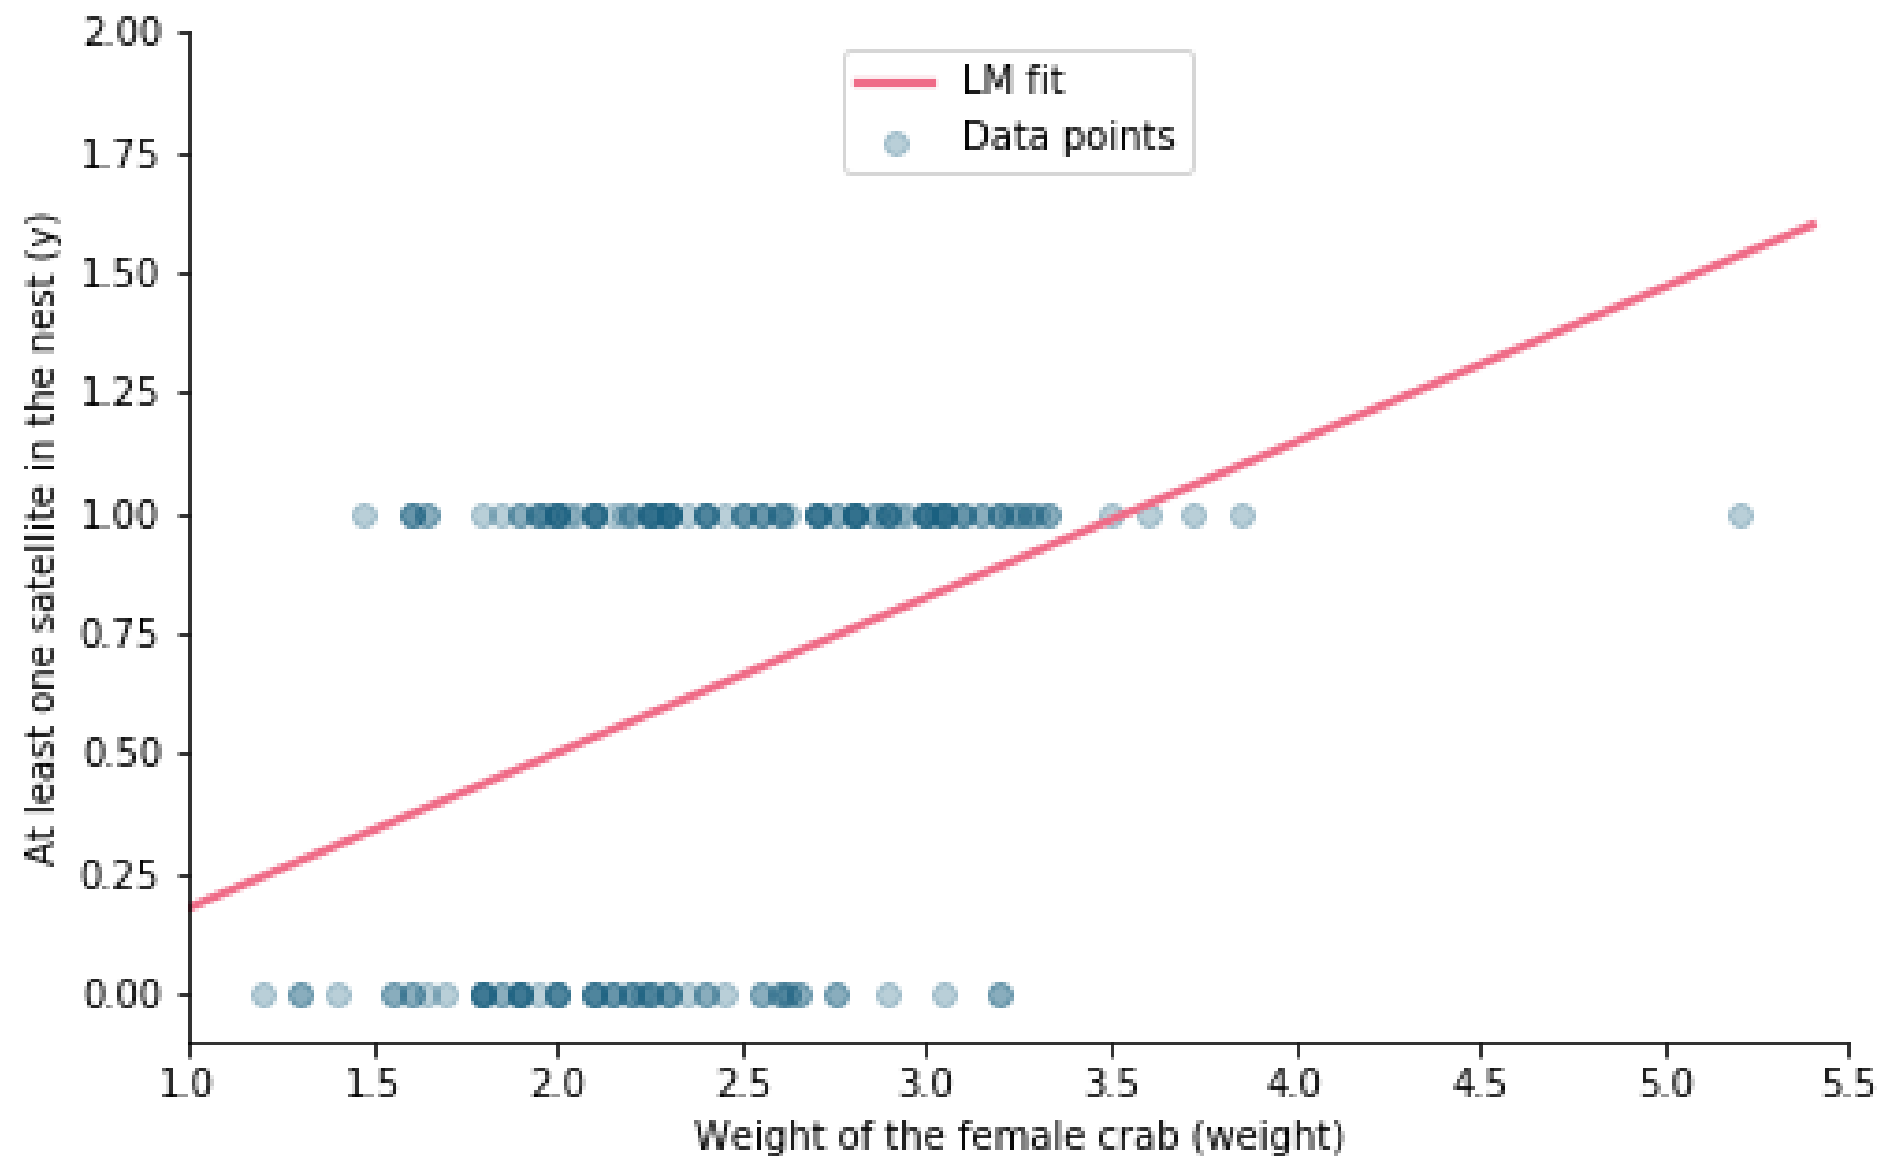
$y \sim \text{weight}$

$$P(\text{satellite crab is present}) = P(y = 1)$$

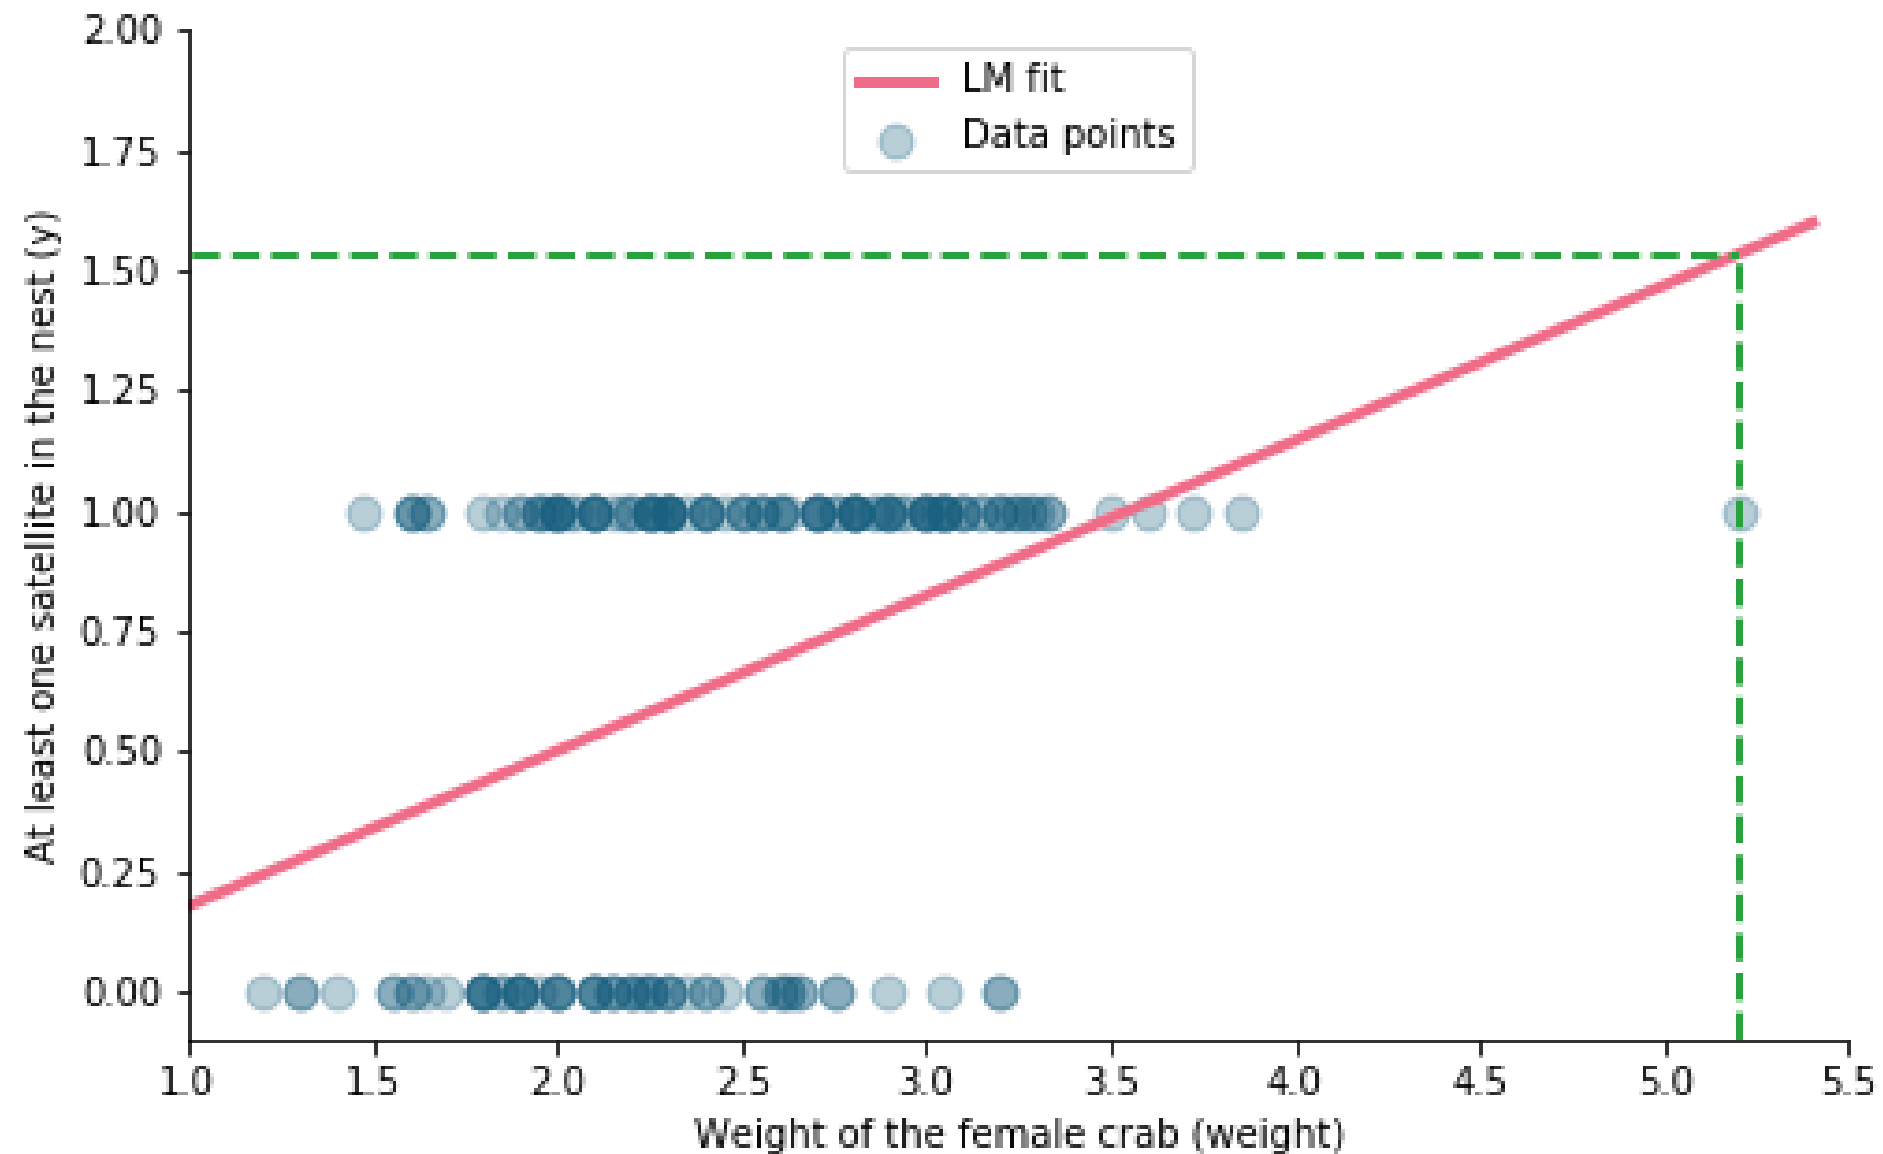
# Linear model and binary response



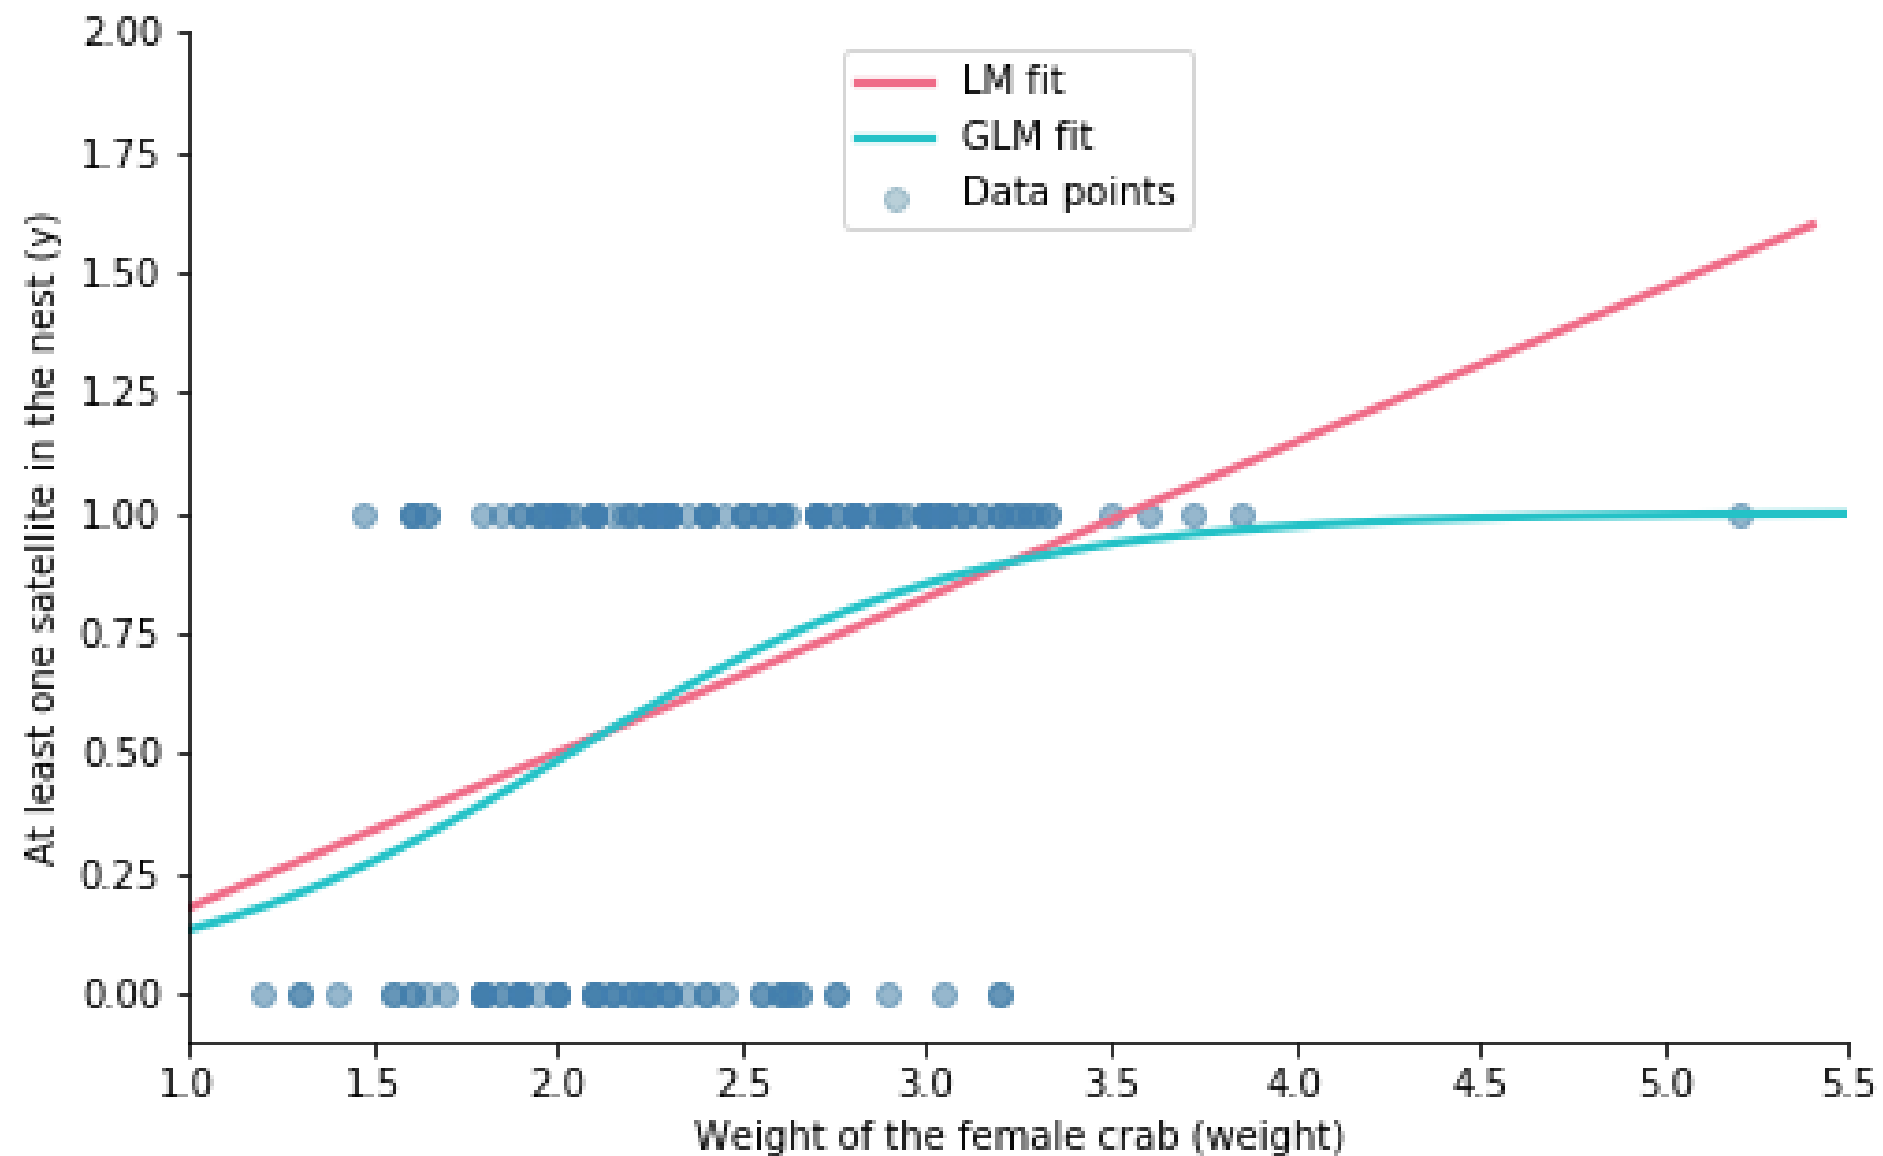
# Linear model and binary response



# Linear model and binary response

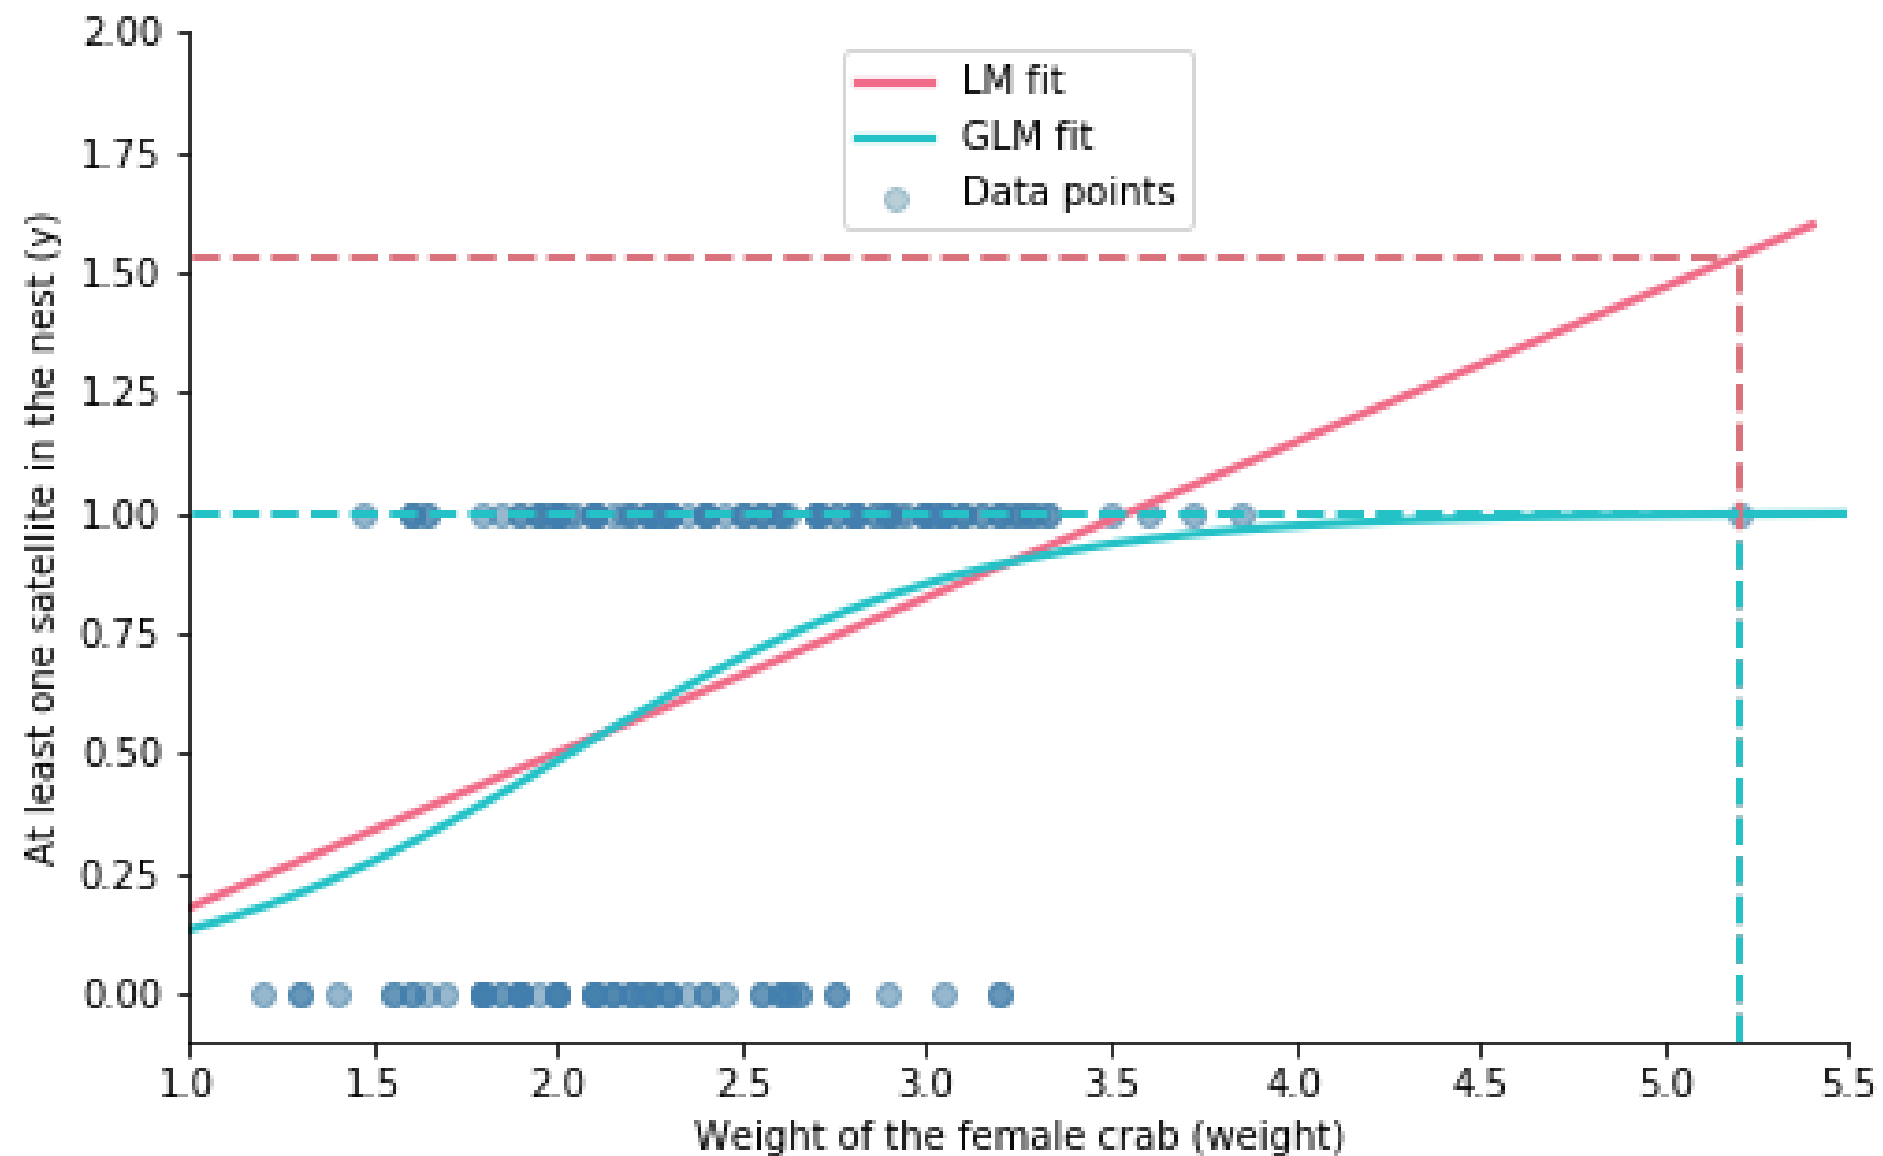


# Linear model and binary data

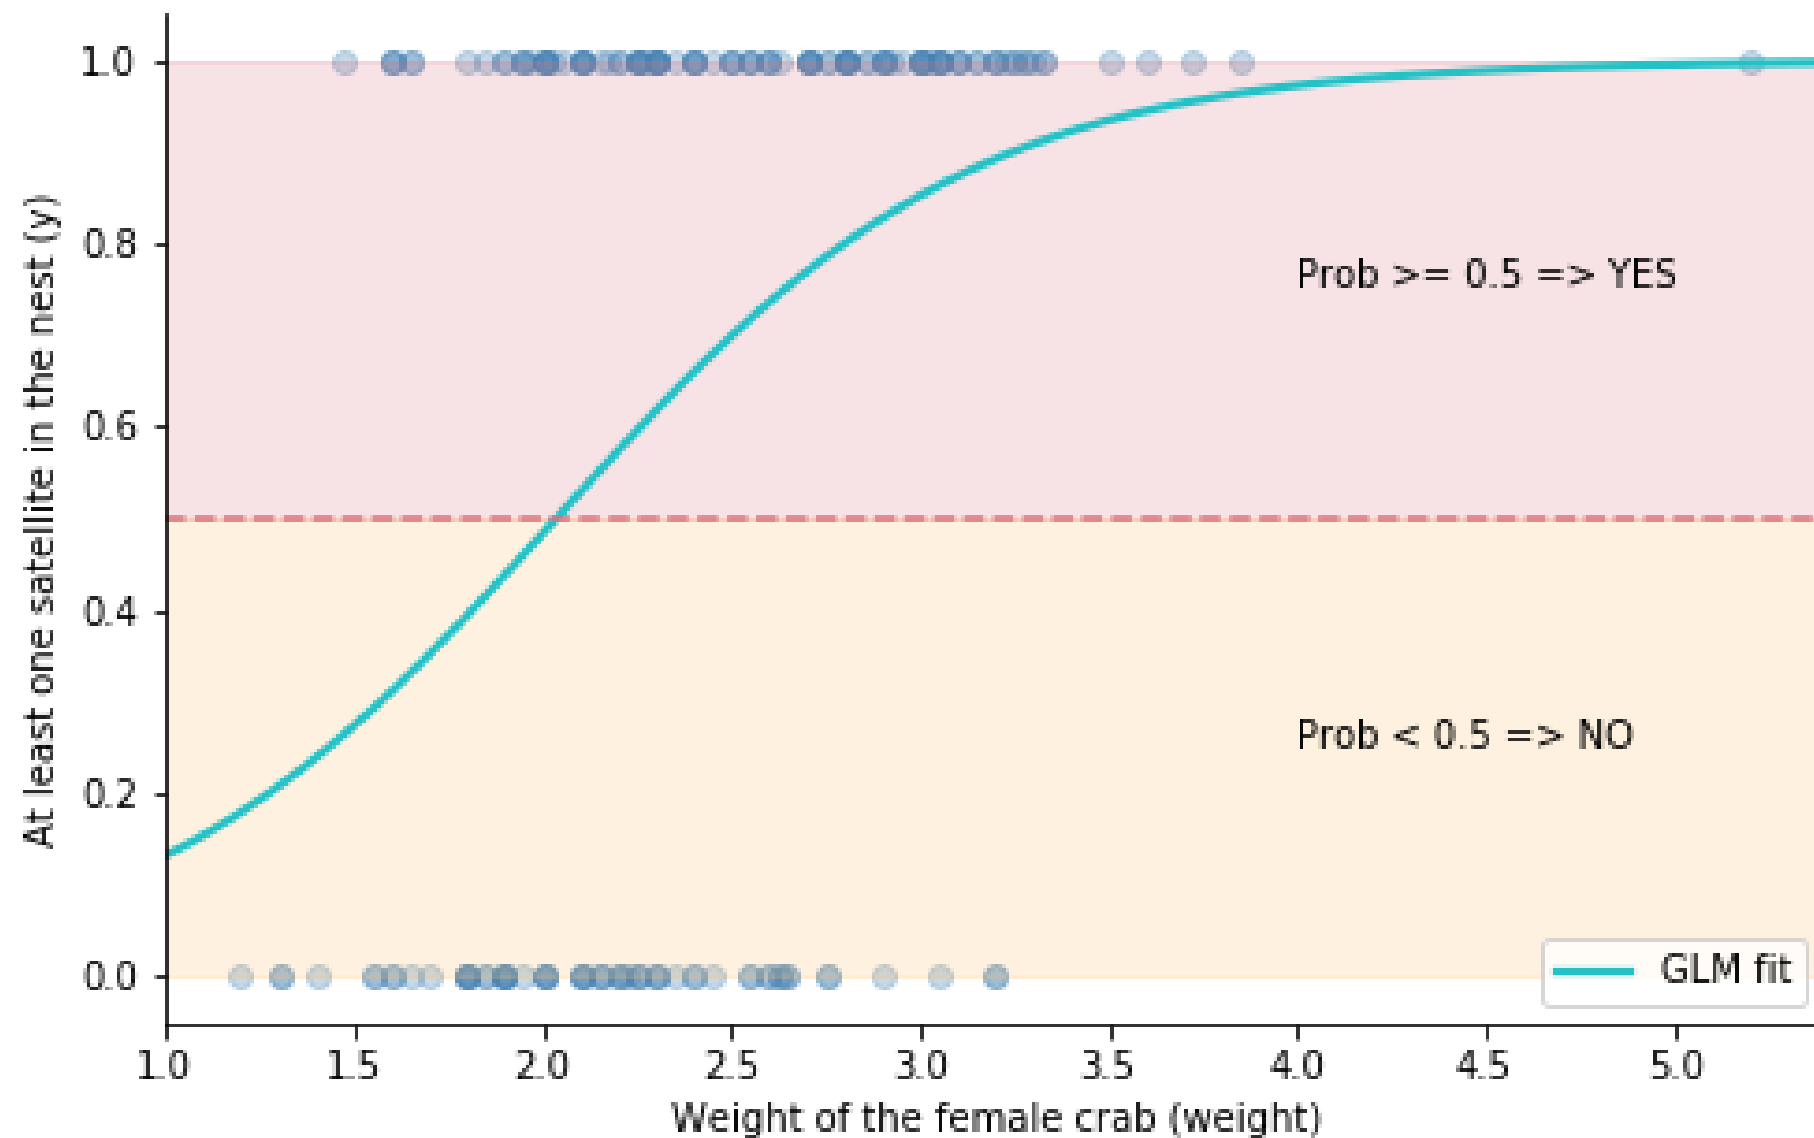




# Linear model and binary data



# From probabilities to classes



# Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

# How to build a GLM?

GENERALIZED LINEAR MODELS IN PYTHON



**Ita Cirovic Donev**  
Data Science Consultant

# Components of the GLM

Random  
Component

$$g(E[y]) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

# Components of the GLM

Random  
Component

$$g(E[\textcolor{red}{y}]) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

Systematic  
Component

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \dots + \beta_p \textcolor{red}{x}_p$$

# Components of the GLM

Random  
Component

$$g(E[\textcolor{red}{y}]) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

Systematic  
Component

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \dots + \beta_p \textcolor{red}{x}_p$$

Interaction

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \beta_1 \textcolor{red}{x}_2 + \beta_3 \textcolor{red}{x}_1 * \textcolor{red}{x}_2$$

# Components of the GLM

Random  
Component

$$g(E[\textcolor{red}{y}]) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

Systematic  
Component

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \dots + \beta_p \textcolor{red}{x}_p$$

Interaction

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \beta_1 \textcolor{red}{x}_2 + \beta_3 \textcolor{red}{x}_1 * \textcolor{red}{x}_2$$

Curvilinear

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \beta_2 \textcolor{red}{x}_1^2$$



# Components of the GLM

Random  
Component

Systematic  
Component


Link  
Function

$$g(E[\textcolor{red}{y}]) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

$$g(E[y]) = \beta_0 + \beta_1 \textcolor{red}{x}_1 + \dots + \beta_p \textcolor{red}{x}_p$$

$$\textcolor{red}{g}(E[y]) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

# Continuous → Linear Regression



Data type: continuous

Domain:  $(-\infty, \infty)$

Examples: house price, salary, person's height

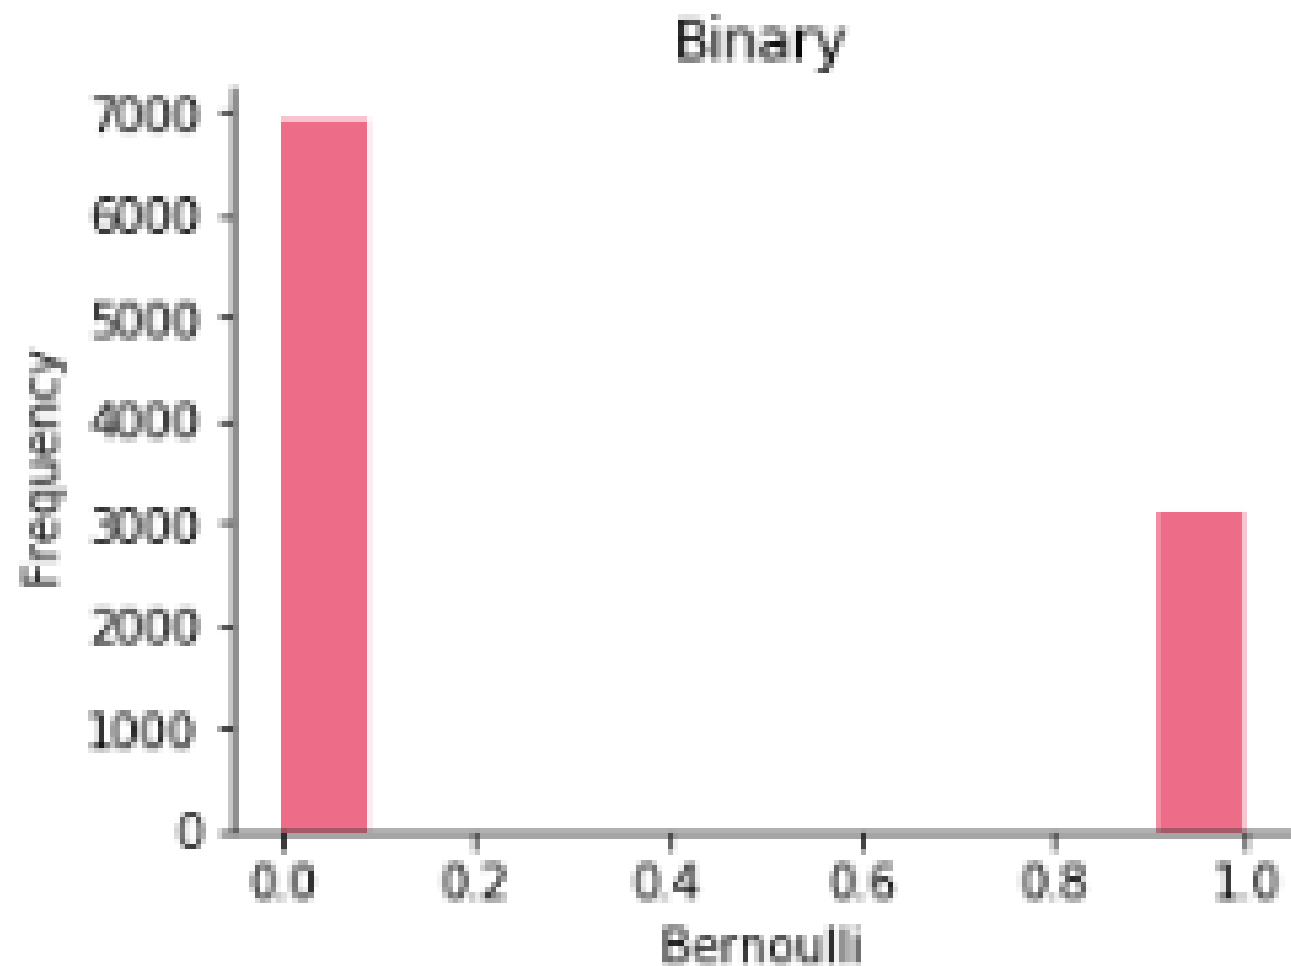
Family: `Gaussian()`

Link: identity

$$g(\mu) = \mu = E(y)$$

Model = Linear regression

# Binary → Logistic regression



Data type: binary

Domain: 0, 1

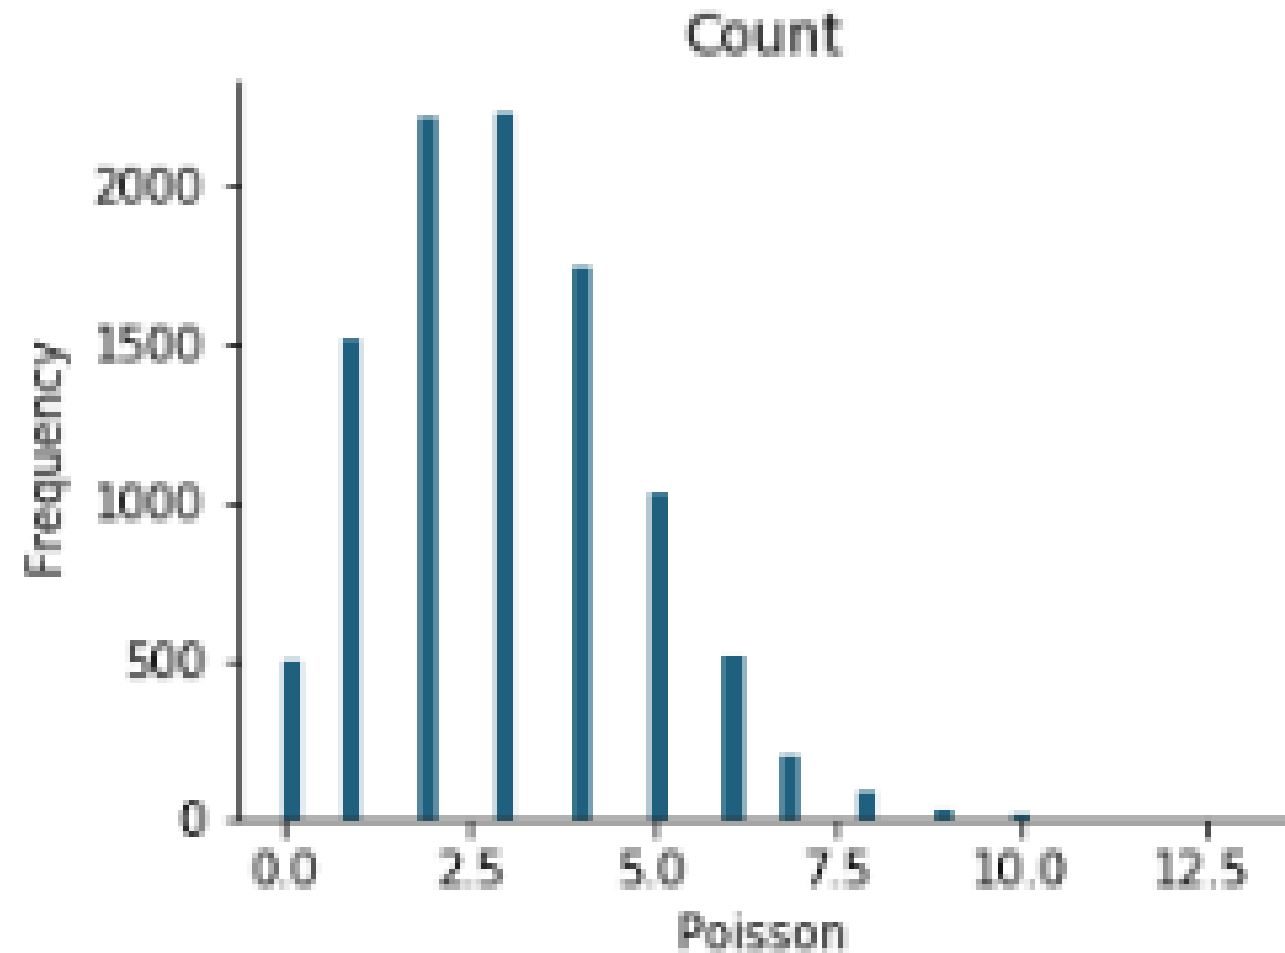
Examples: True/False

Family: `Binomial()`

Link: logit

Model = Logistic regression

# Count → Poisson regression



Data type: count

Domain:  $0, 1, 2, \dots, \infty$

Examples: number of votes, number of hurricanes

Family: `Poisson()`

Link: logarithm

Model = Poisson regression

# Link functions

Density	Link: $\eta = g(\mu)$	Default link	<code>glm(family=...)</code>
Normal	$\eta = \mu$	identity	<code>Gaussian()</code>
Poisson	$\eta = \log(\mu)$	logarithm	<code>Poisson()</code>
Binomial	$\eta = \log[p/(1 - p)]$	logit	<code>Binomial()</code>
Gamma	$\eta = 1/\mu$	inverse	<code>Gamma()</code>
Inverse Gaussian	$\eta = 1/\mu^2$	inverse squared	<code>InverseGaussian()</code>

# Benefits of GLMs

- A unified framework for many different data distributions
  - Exponential family of distributions
- Link function
  - Transforms the expected value of  $y$
  - Enables linear combinations
  - Many techniques from linear models apply to GLMs as well

# Let's practice

GENERALIZED LINEAR MODELS IN PYTHON

# How to fit a GLM in Python?

GENERALIZED LINEAR MODELS IN PYTHON



**Ita Cirovic Donev**  
Data Science Consultant



# statsmodels

- Importing statsmodels

```
import statsmodels.api as sm
```

- Support for formulas

```
import statsmodels.formula.api as smf
```

- Use `glm()` directly

```
from statsmodels.formula.api import glm
```

# Process of model fit

1. Describe the model → `glm()`
2. Fit the model → `.fit()`
3. Summarize the model → `.summary()`
4. Make model predictions → `.predict()`

# Describing the model

## FORMULA based

```
from statsmodels.formula.api import glm
```

```
model = glm(formula, data, family)
```

## ARRAY based

```
import statsmodels.api as sm
```

```
X = sm.add_constant(X)  
model = sm.glm(y, X, family)
```

# Formula Argument

`response ~ explanatory variable(s)`

`output ~ input(s)`

```
formula = 'y ~ x1 + x2'
```

- `C(x1)` : treat `x1` as categorical variable
- `-1` : remove intercept
- `x1:x2` : an interaction term between `x1` and `x2`
- `x1*x2` : an interaction term between `x1` and `x2` and the individual variables
- `np.log(x1)` : apply vectorized functions to model variables

# Family Argument

```
family = sm.families.____()
```

The family functions:

- `Gaussian(link = sm.families.links.identity)` → the default family
- `Binomial(link = sm.families.links.logit)`
  - `probit`, `cauchy`, `log`, and `cloglog`
- `Poisson(link = sm.families.links.log)`
  - `identity` and `sqrt`

Other distribution families you can review at [statsmodels website](#).

# Summarizing the model

```
print(model_GLM.summary())
```

## Generalized Linear Model Regression Results

```
=====
Dep. Variable:          y   No. Observations:          173
Model:                GLM   Df Residuals:            171
Model Family:        Binomial   Df Model:              1
Link Function:        logit   Scale:                1.0000
Method:              IRLS   Log-Likelihood:        -97.226
Date:                Mon, 21 Jan 2019   Deviance:          194.45
Time:                11:30:01   Pearson chi2:          165.
No. Iterations:        4   Covariance Type:            nonrobust
=====
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-12.3508	2.629	-4.698	0.000	-17.503	-7.199
width	0.4972	0.102	4.887	0.000	0.298	0.697

# Regression coefficients

`.params` prints regression coefficients

```
model_GLM.params
```

```
Intercept    -12.350818  
width         0.497231  
dtype: float64
```

`.conf_int(alpha=0.05, cols=None)`

prints confidence intervals

```
model_GLM.conf_int()
```

```
              0          1  
Intercept -17.503010 -7.198625  
width      0.297833  0.696629
```



# Predictions

- Specify all the model variables in test data
- `.predict(test_data)` computes predictions

```
model_GLM.predict(test_data)
```

```
0    0.029309
1    0.470299
2    0.834983
3    0.972363
4    0.987941
```

# Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON