

# Binary data and logistic regression

GENERALIZED LINEAR MODELS IN PYTHON



**Ita Cirovic Donev**  
Data Science Consultant

# Binary response data

- Two-class response  $\rightarrow 0, 1$

Examples:

- Credit scoring  $\rightarrow$  "Default" / "Non-Default"
- Passing a test  $\rightarrow$  "Pass" / "Fail"
- Fraud detection  $\rightarrow$  "Fraud" / "No-Fraud"
- Choice of a product  $\rightarrow$  "Product ABC" / "Product XYZ"

# Binary data

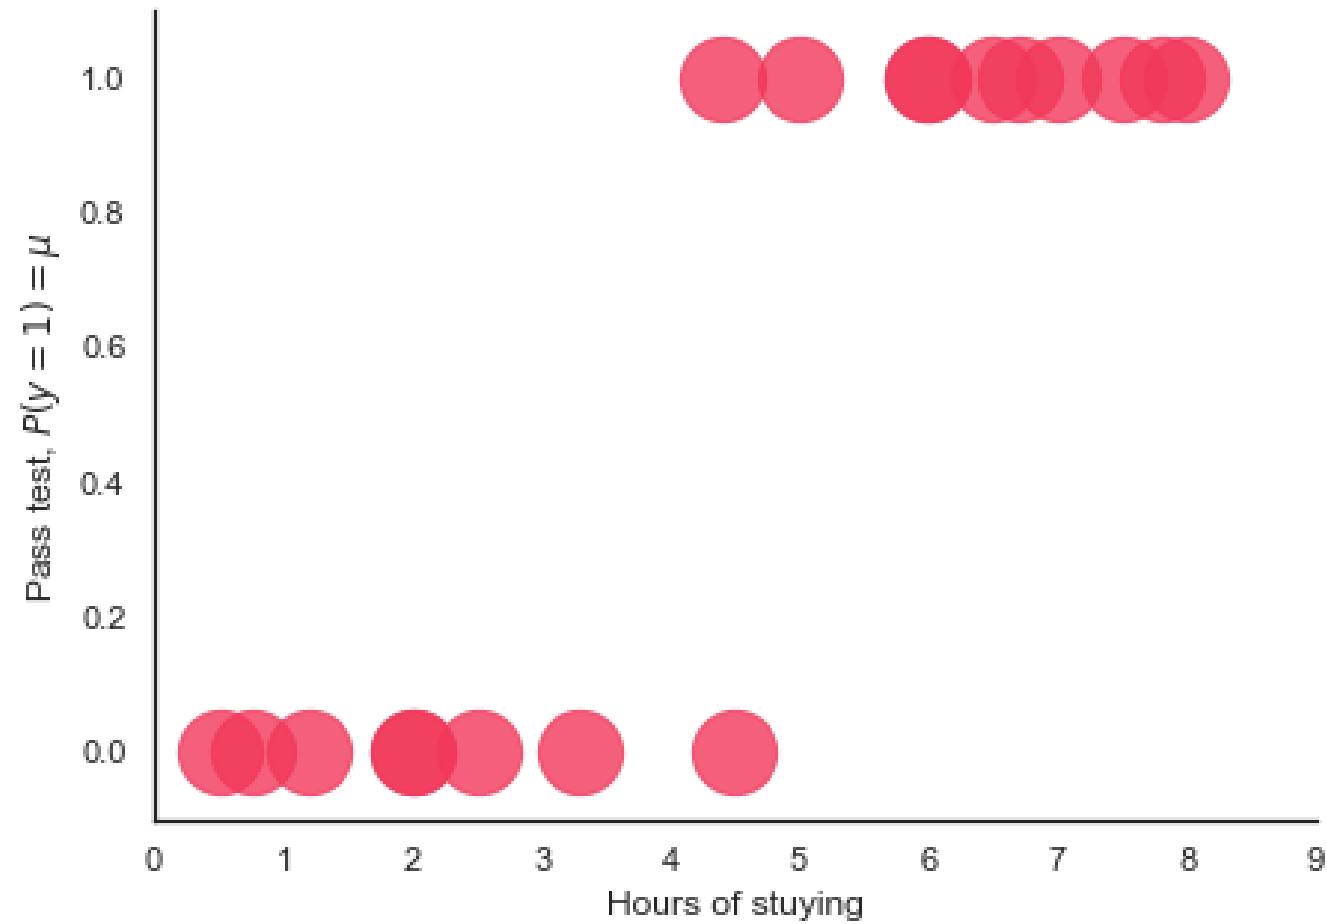
## UNGROUPED

- Single event
- Flip one coin
- Two of possible outcomes: 0/1
- $Bernoulli(p)$  or
- $Binomial(n = 1, p)$

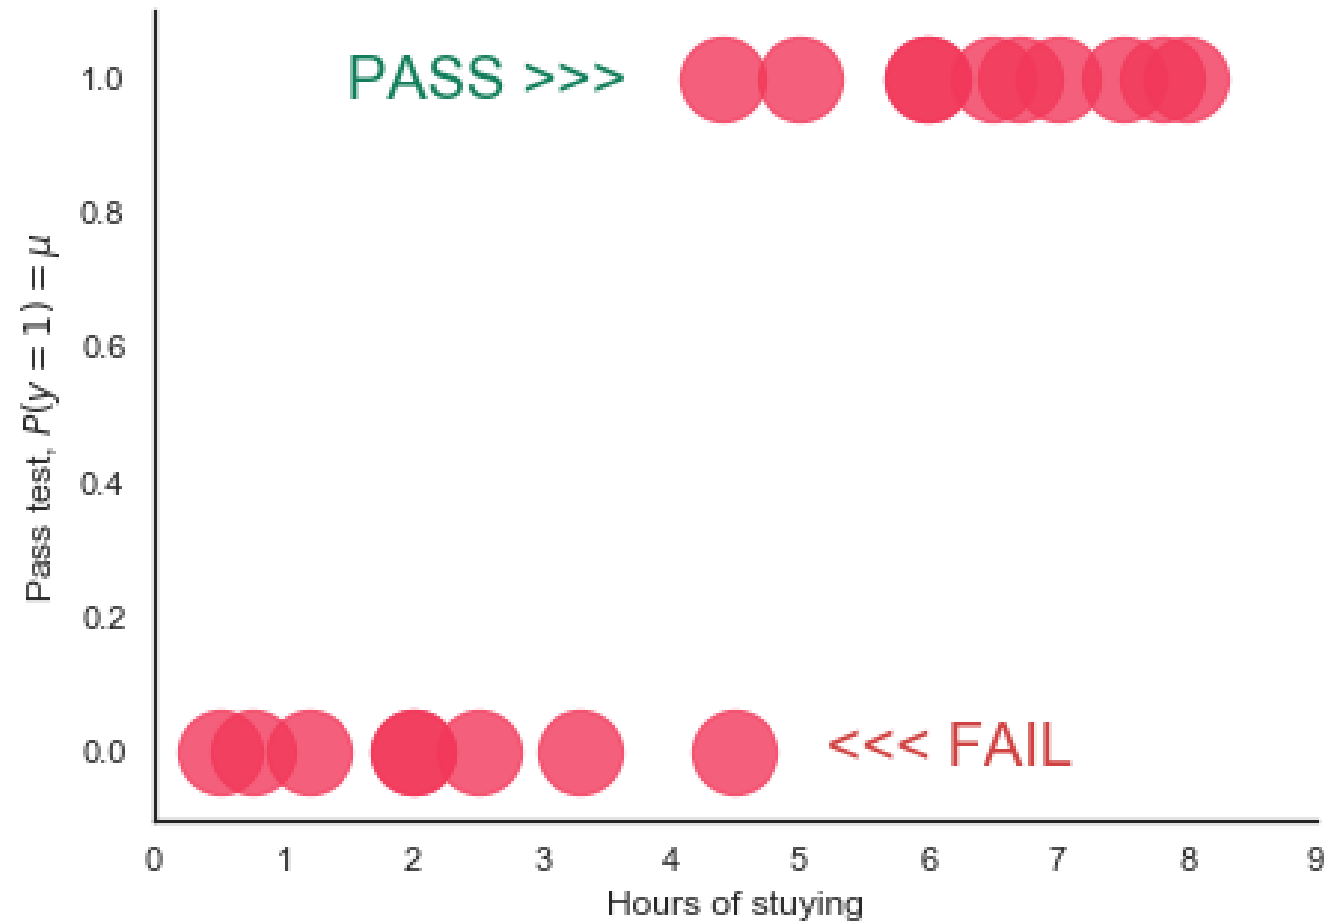
## GROUPED

- Multiple events
- Flip multiple coins
- Number of successes in a given  $n$  number of trials
- $Binomial(n, p)$

# Logistic function



# Logistic function

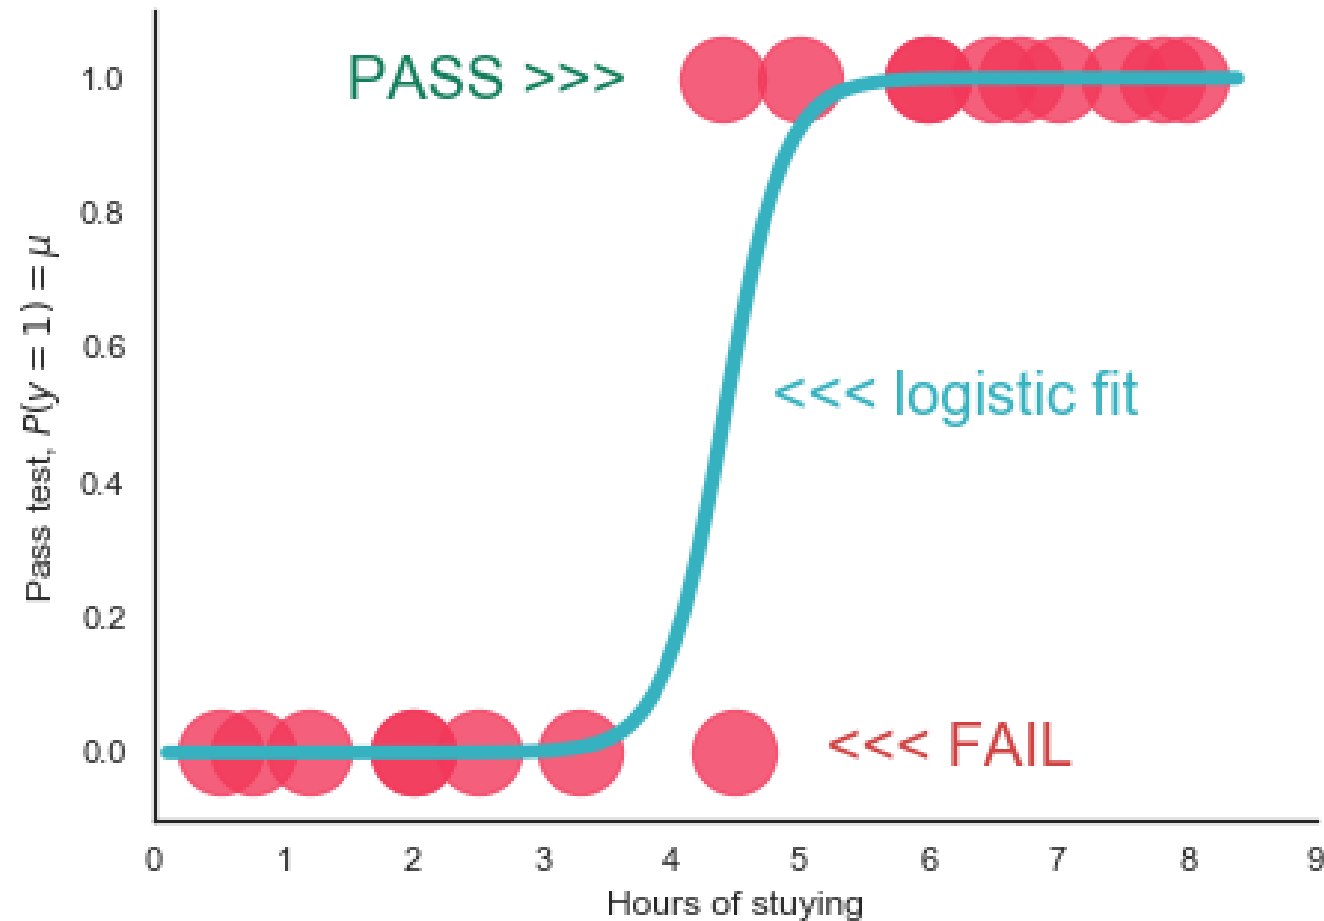


- Test outcome:  $PASS = 1$  or  $FAIL = 0$
- Want to model

$$P(y = 1) = \beta_0 + \beta_1 x_1$$

$$P(\text{Pass}) = \beta_0 + \beta_1 \times \text{Hours of study}$$

# Logistic function



- Test outcome:  $PASS = 1$  or  $FAIL = 0$
- Want to model

$$P(y = 1) = \beta_0 + \beta_1 x_1$$

$$P(\text{Pass}) = \beta_0 + \beta_1 \times \text{Hours of study}$$

- Use logistic function

$$f(z) = \frac{1}{(1 + \exp(-z))}$$

# Odds and odds ratio

$$ODDS = \frac{\text{event occurring}}{\text{event NOT occurring}}$$

$$ODDS\ RATIO = \frac{odds1}{odds2}$$

# Odds example

- 4 games

Win Win Win Lose

- Odds are 3 to 1

$$\text{Odds} = \frac{\text{Win Win Win}}{\text{Lose}}$$



# Odds and probabilities

odds  $\neq$  probability

$$\text{odds} = \frac{\text{probability}}{1 - \text{probability}}$$

$$\text{probability} = \frac{\text{odds}}{1 + \text{odds}}$$

# From probability model to logistic regression

## Step 1. Probability model

$$E(y) = \mu = P(y = 1) = \beta_0 + \beta_1 x_1$$

## Step 2. Logistic function

$$f(z) = \frac{1}{(1 + \exp(-z))}$$

## Step 3. Apply logistic function → INVERSE-LOGIT

$$\mu = \frac{1}{1 + \exp(-(\beta_0 + \beta_1 x_1))} = \frac{\exp(\beta_0 + \beta_1 x_1)}{1 + \exp(\beta_0 + \beta_1 x_1)}$$

$$1 - \mu = \frac{1}{1 + \exp(\beta_0 + \beta_1 x_1)}$$

# From probability model to logistic regression

- Probability  $\rightarrow$  odds

$$ODDS = \frac{\mu}{1 - \mu} = \exp(\beta_0 + \beta_1 x_1)$$

- Log transformation  $\rightarrow$  **LOGISTIC REGRESSION**

$$LOGIT(\mu) = \log\left(\frac{\mu}{1 - \mu}\right) = \beta_0 + \beta_1 x_1$$

# Logistic regression in Python

Function - `glm()`

```
model_GLM = glm(formula = 'y ~ x',  
                 data = my_data,  
                 family = sm.families.Binomial())a).fit
```

Input

```
y = [0, 1, 1, 0, ...]  
y = ['No', 'Yes', 'Yes', ...]  
y = ['Fail', 'Pass', 'Pass', ...]
```

# Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

# Interpreting coefficients

GENERALIZED LINEAR MODELS IN PYTHON



**Ita Cirovic Donev**  
Data Science Consultant

# Model coefficients

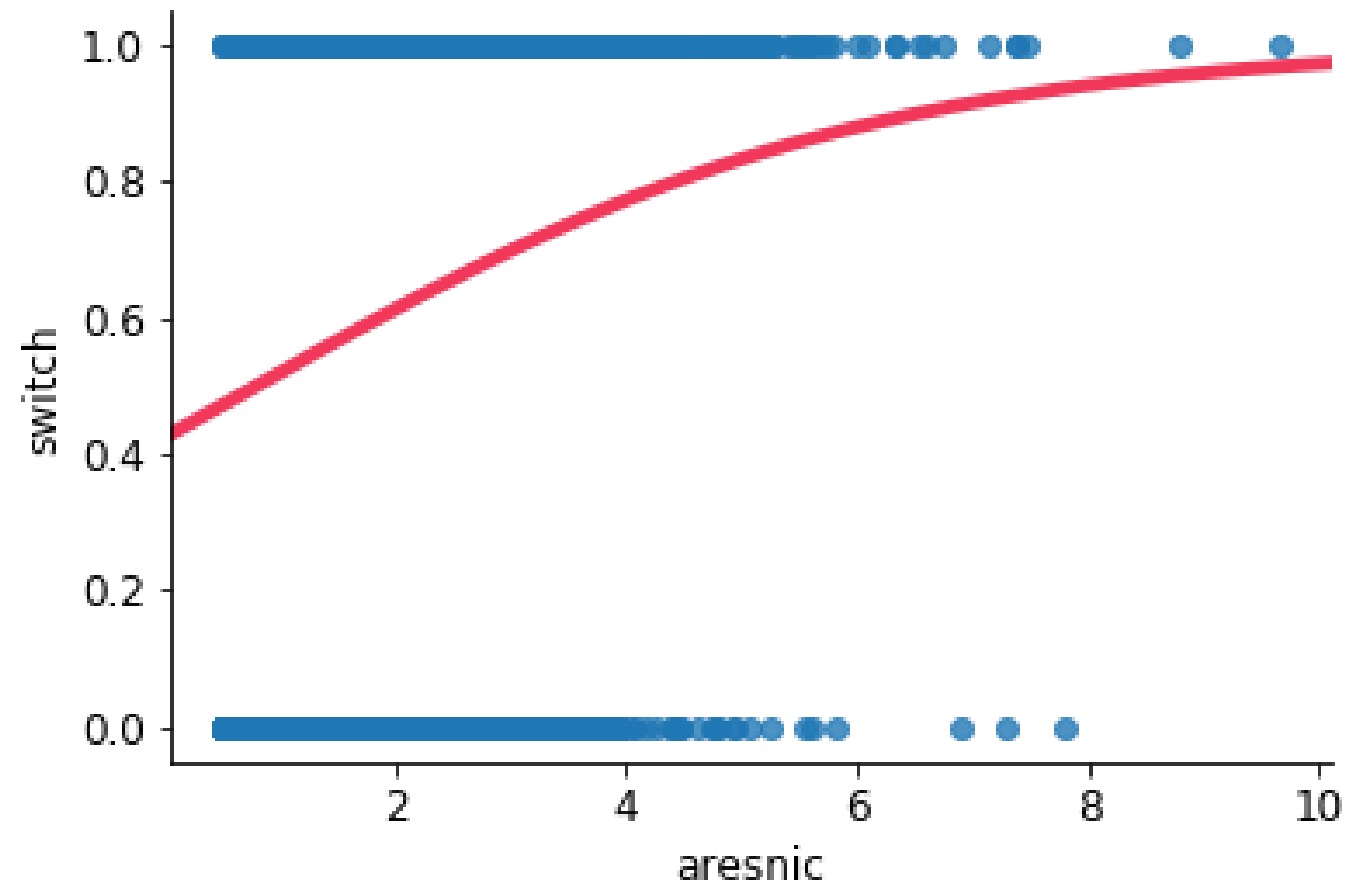
## Generalized Linear Model Regression Results

```
=====
Dep. Variable:                y      No. Observations:          173
Model:                      GLM      Df Residuals:             171
Model Family:              Binomial  Df Model:                 1
Link Function:              logit    Scale:                   1.0000
Method:                     IRLS     Log-Likelihood:          -97.869
Date:                       Sat, 23 Feb 2019    Deviance:                195.74
Time:                       13:03:32    Pearson chi2:            168.
No. Iterations:              5      Covariance Type:         nonrobust
=====
```

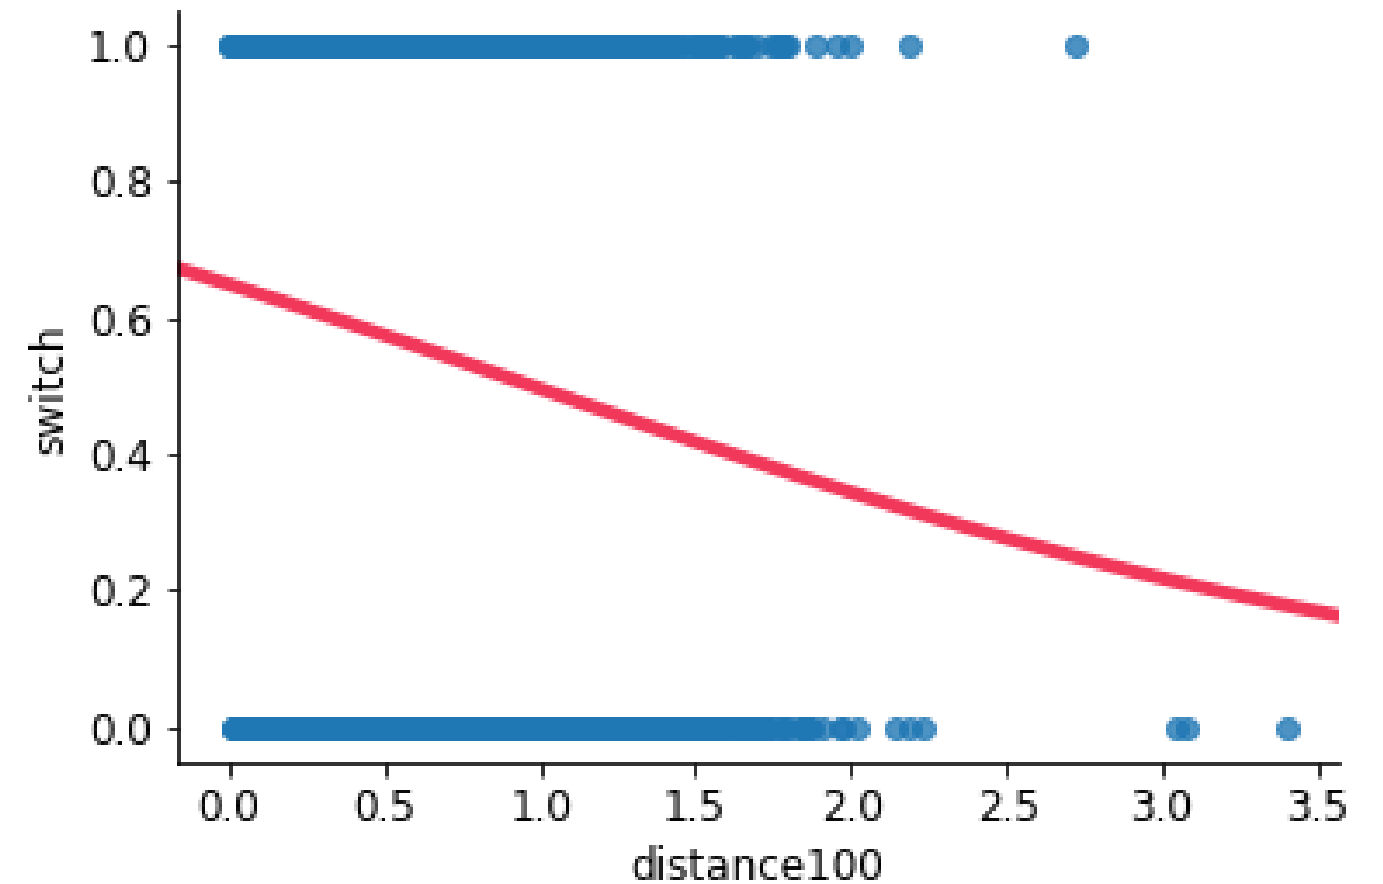
	coef	std err	z	P> z	[0.025	0.975]
Intercept	-3.6947	0.880	-4.198	0.000	-5.420	-1.970
weight	1.8151	0.377	4.819	0.000	1.077	2.553

# Coefficient beta

- $\beta > 0 \rightarrow$  ascending curve



- $\beta < 0 \rightarrow$  descending curve





# Linear vs logistic

## LINEAR MODEL

```
glm('y ~ weight',  
    data = crab,  
    family = sm.families.Gaussian())
```

$$\mu = -0.14 + 0.32 * weight$$

For every *one-unit* increase in weight

- **estimated probability** increases by 0.32

## LOGIT MODEL

```
glm('y ~ weight',  
    data = crab,  
    family = sm.families.Binomial())
```

$$\log(odds) = -3.69 + 1.8 * weight$$

For every *one-unit* increase in weight

- **log(odds)** increase by 1.8

# Log odds interpretation

- Logistic model

$$\log\left(\frac{\mu}{1-\mu}\right) = \beta_0 + \beta_1 x_1$$

- Increase  $x$  by one-unit

$$\log\left(\frac{\mu}{1-\mu}\right) = \beta_0 + \beta_1 (x_1 + 1)$$

# Log odds interpretation

- Logistic model

$$\log\left(\frac{\mu}{1-\mu}\right) = \beta_0 + \beta_1 x_1$$

- Increase  $x$  by one-unit

$$\log\left(\frac{\mu}{1-\mu}\right) = \beta_0 + \beta_1 (x_1 + 1) = \beta_0 + \beta_1 x_1 + \beta_1$$

- Take the exponential

$$\left(\frac{\mu}{1-\mu}\right) = \exp(\beta_0 + \beta_1 x_1) \exp(\beta_1)$$

**Conclusion** → the **odds** are multiplied by  $\exp(\beta_1)$

# Log odds interpretation

- Crab model `y ~ weight`

$$\log\left(\frac{\mu}{1-\mu}\right) = -3.6947 + 1.815 * weight$$

- The odds of satellite crab multiply by  $\exp(1.815) = 6.14$  for a unit increase in weight

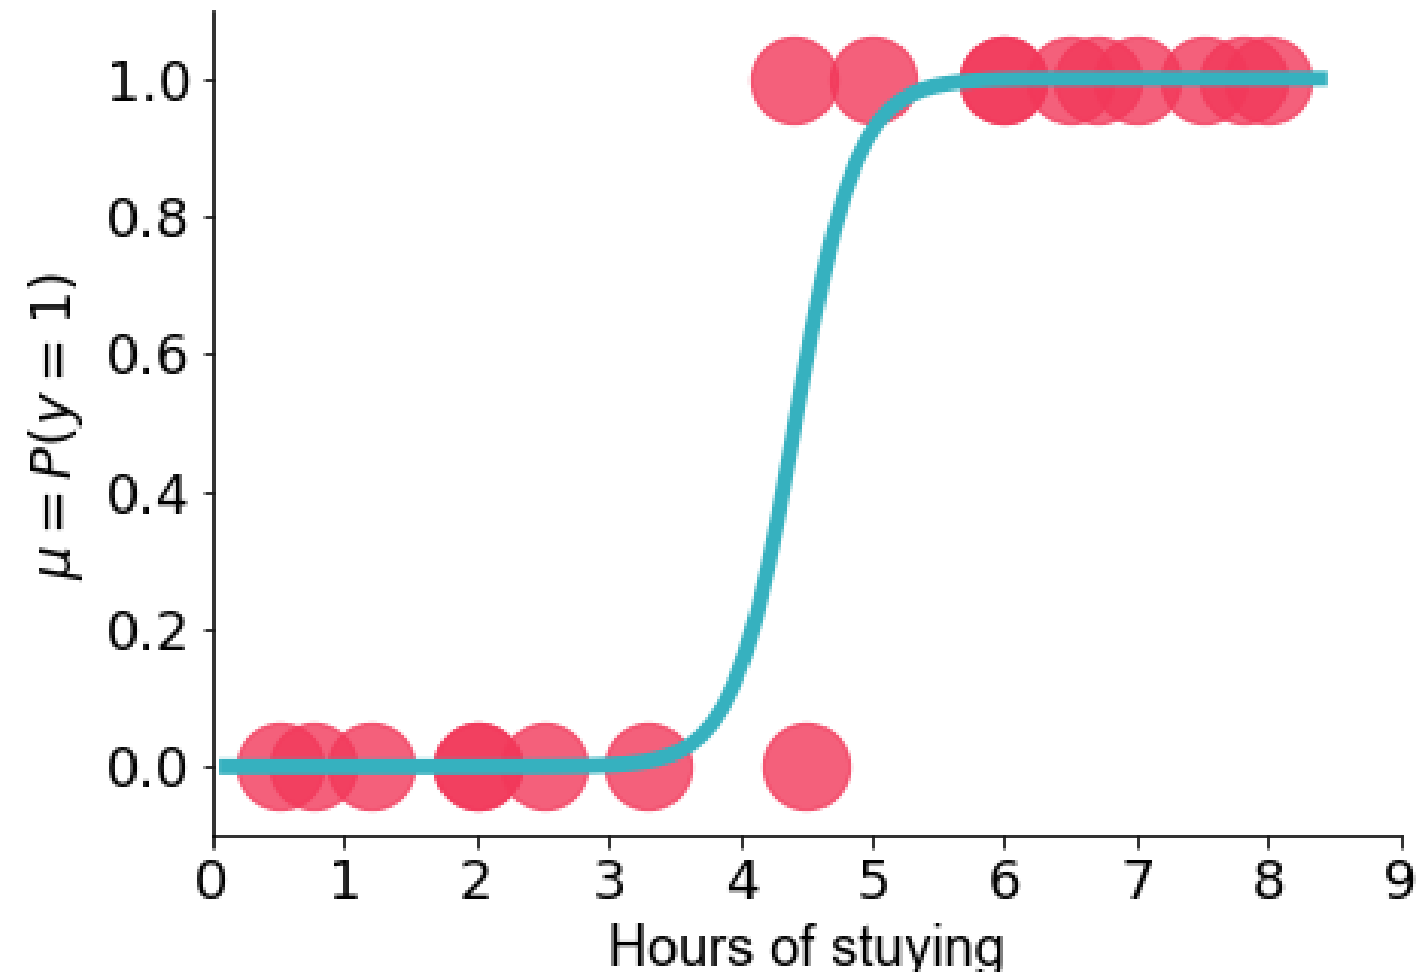
# Log odds interpretation

- Crab model  $y \sim \text{weight}$


$$\log\left(\frac{\mu}{1 - \mu}\right) = -3.6947 + 1.8151 * \text{weight}$$

- The odds of satellite crab multiply by  $\exp(1.8151) = 6.14$  for a unit increase in weight
- The intercept coefficient of  $-3.6947$  denotes the **baseline log odds**
  - $\exp(-3.6947) = 0.0248$  are the odds when  $\text{weight} = 0$ .

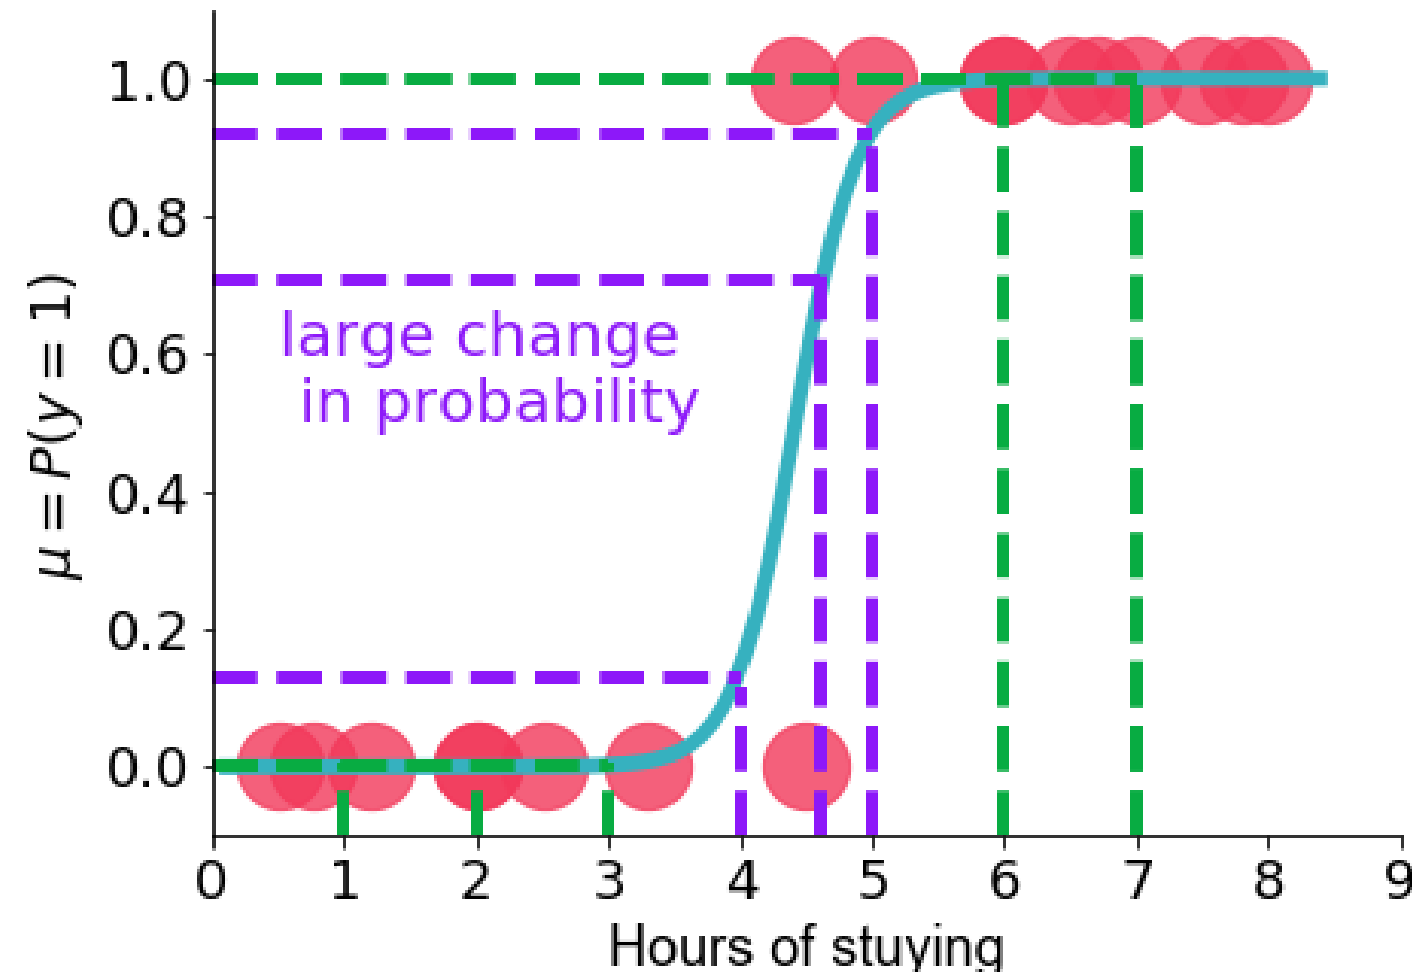
# Probability vs logistic fit



# Probability vs logistic fit



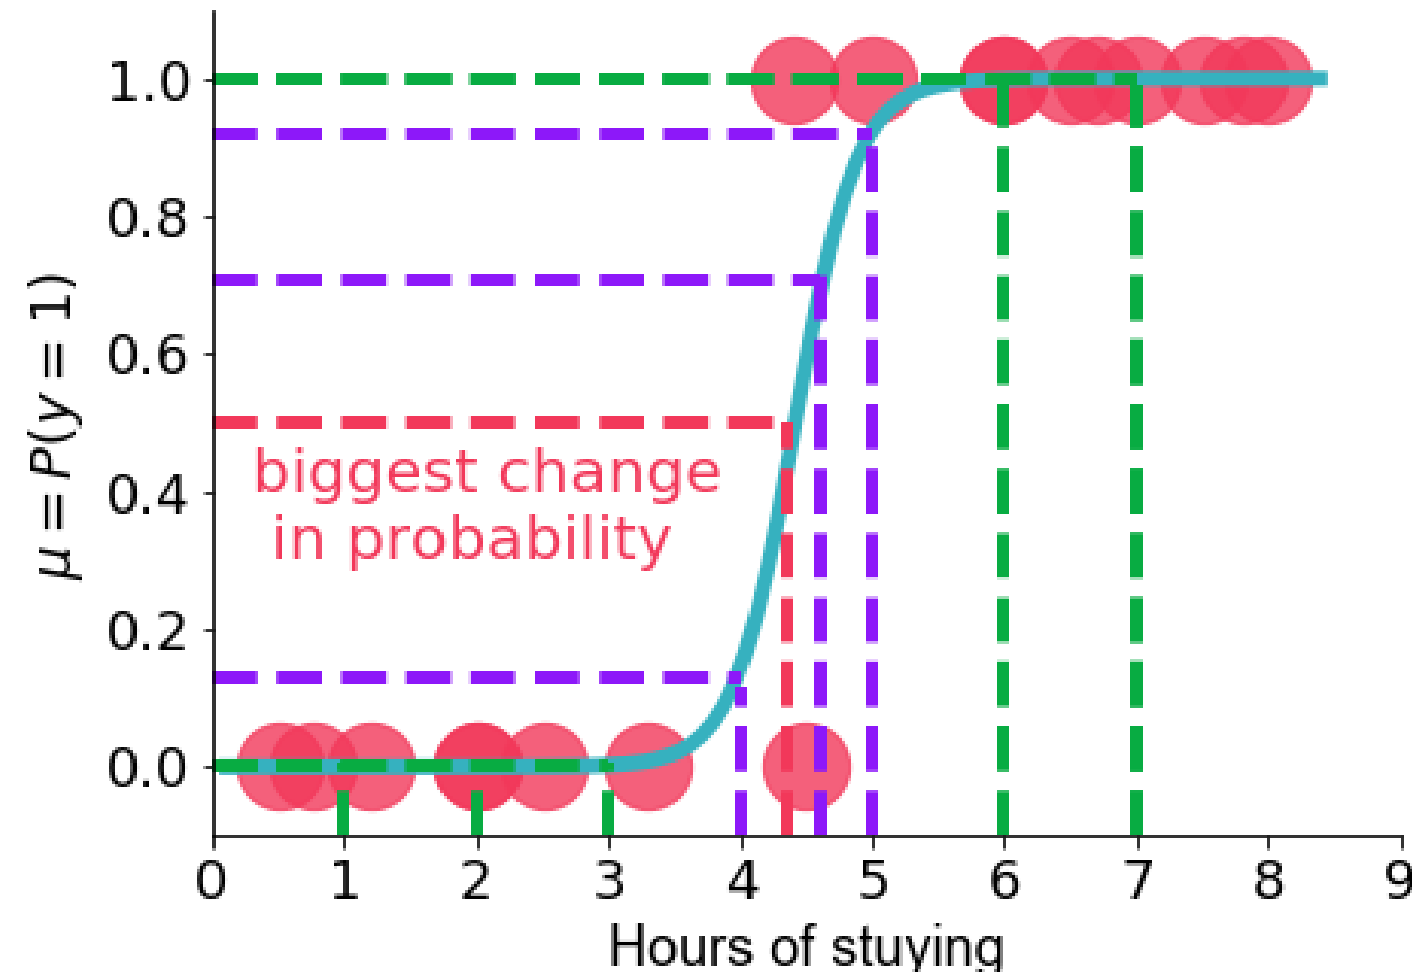
# Probability vs logistic fit



- slope  $\rightarrow \beta \times \mu(1 - \mu)$



# Probability vs logistic fit



- slope  $\rightarrow \beta \times \mu(1 - \mu)$

# Compute change in estimated probability

```
# Choose x (weight) and extract model coefficients
x = 1.5
intercept, slope = model_GLM.params
```

```
# Compute estimated probability
est_prob = np.exp(intercept + slope * x) / (1 + np.exp(intercept + slope * x))
```


```
0.2744
```

```
# Compute incremental change in estimated probability given x
ic_prob = slope * est_prob * (1 - est_prob)
```

```
0.3614
```

# Rate of change in probability for every x

$$\text{logit} = -3.6947 + 1.8151 * \text{weight}$$

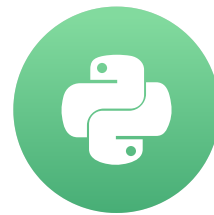


# Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

# Interpreting model inference

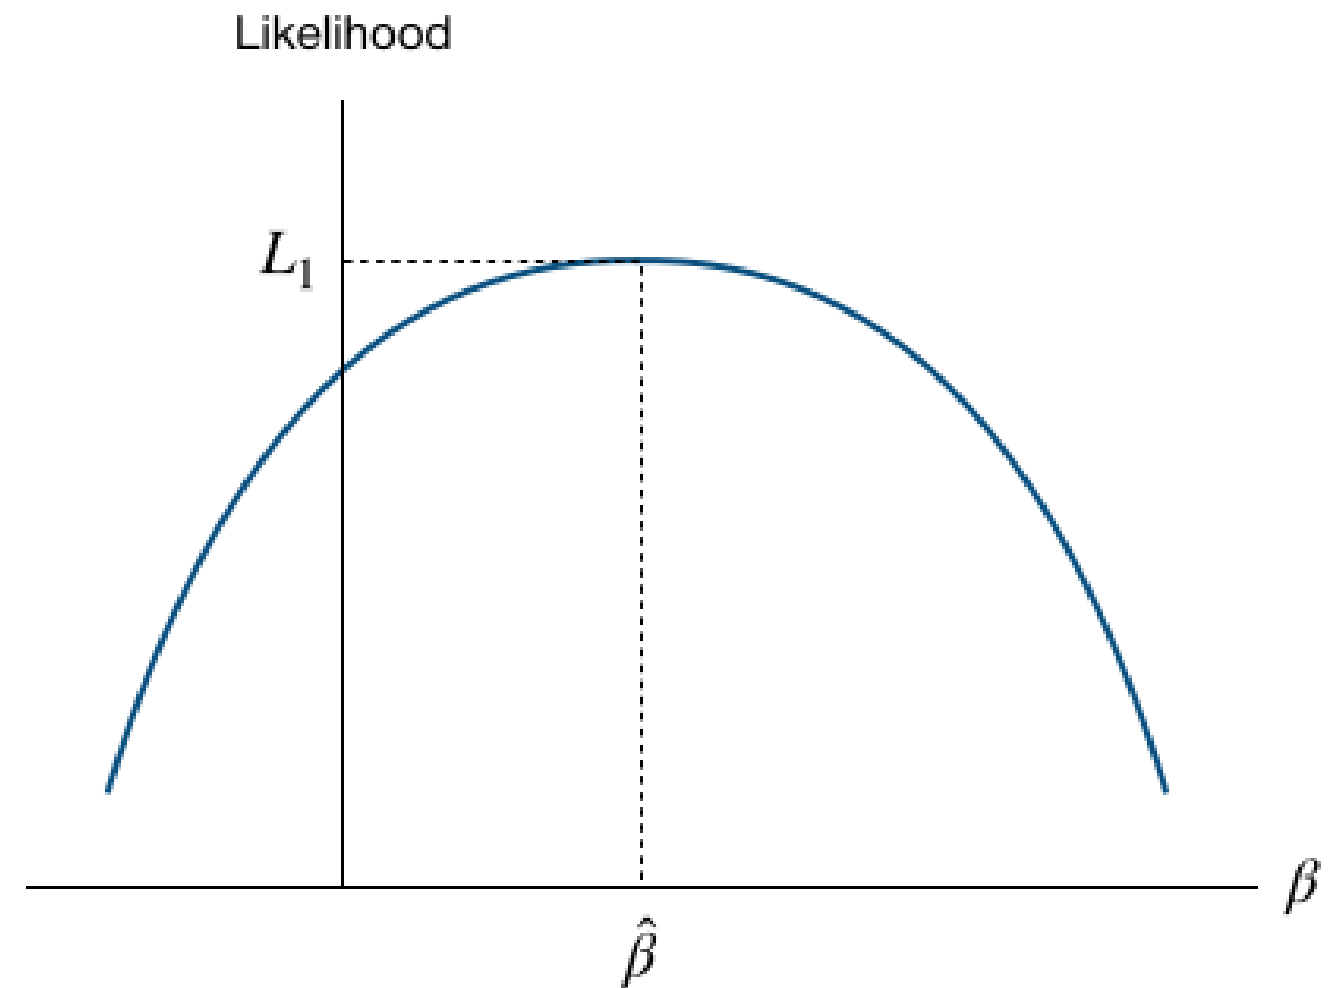
GENERALIZED LINEAR MODELS IN PYTHON



**Ita Cirovic Donev**  
Data Science Consultant

# Estimation of beta coefficient

- Maximum likelihood estimation (MLE)
- Estimated coefficient,  $\hat{\beta}$ 
  - log-likelihood takes on the maximum value



# Estimation of beta coefficient

- Iteratively reweighted least squares (IRLS)

```
Generalized Linear Model Regression Results
=====
Dep. Variable:                y      No. Observations:          173
Model:                        GLM      Df Residuals:              171
Model Family:                 Binomial  Df Model:                  1
Link Function:                 logit     Scale:                    1.0000
Method:                        IRLS      Log-Likelihood:           -97.869
Date:                          Sun, 24 Feb 2019    Deviance:                 195.74
Time:                          12:18:44           Pearson chi2:             168.
No. Iterations:                5          Covariance Type:         nonrobust
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
Intercept    -3.6947      0.880     -4.198     0.000     -5.420     -1.970
weight        1.8151      0.377      4.819     0.000      1.077      2.553
=====
```

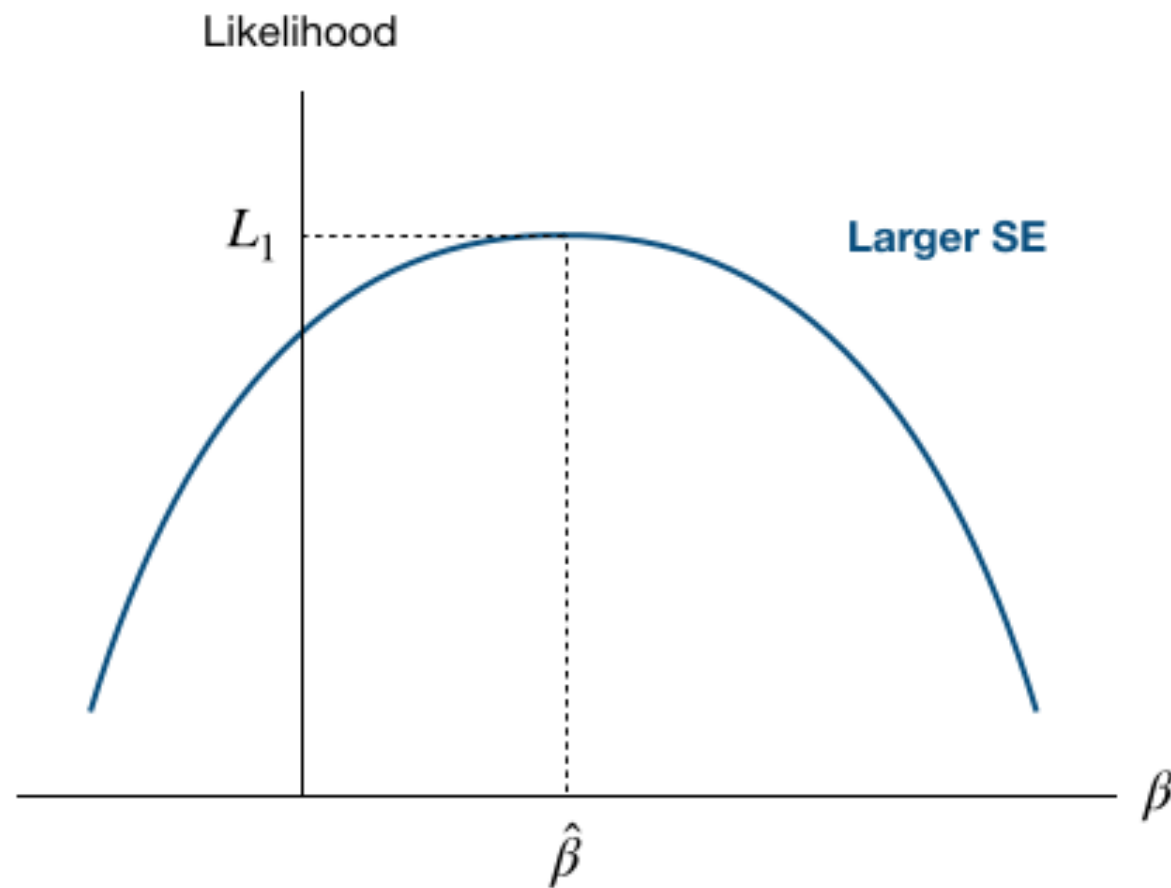
# Significance testing

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	y	No. Observations:	173			
Model:	GLM	Df Residuals:	171			
Model Family:	Binomial	Df Model:	1			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-97.869			
Date:	Sun, 24 Feb 2019	Deviance:	195.74			
Time:	12:18:44	Pearson chi2:	168.			
No. Iterations:	5	Covariance Type:	nonrobust			
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
Intercept	-3.6947	0.880	-4.198	0.000	-5.420	-1.970
weight	1.8151	0.377	4.819	0.000	1.077	2.553
=====						

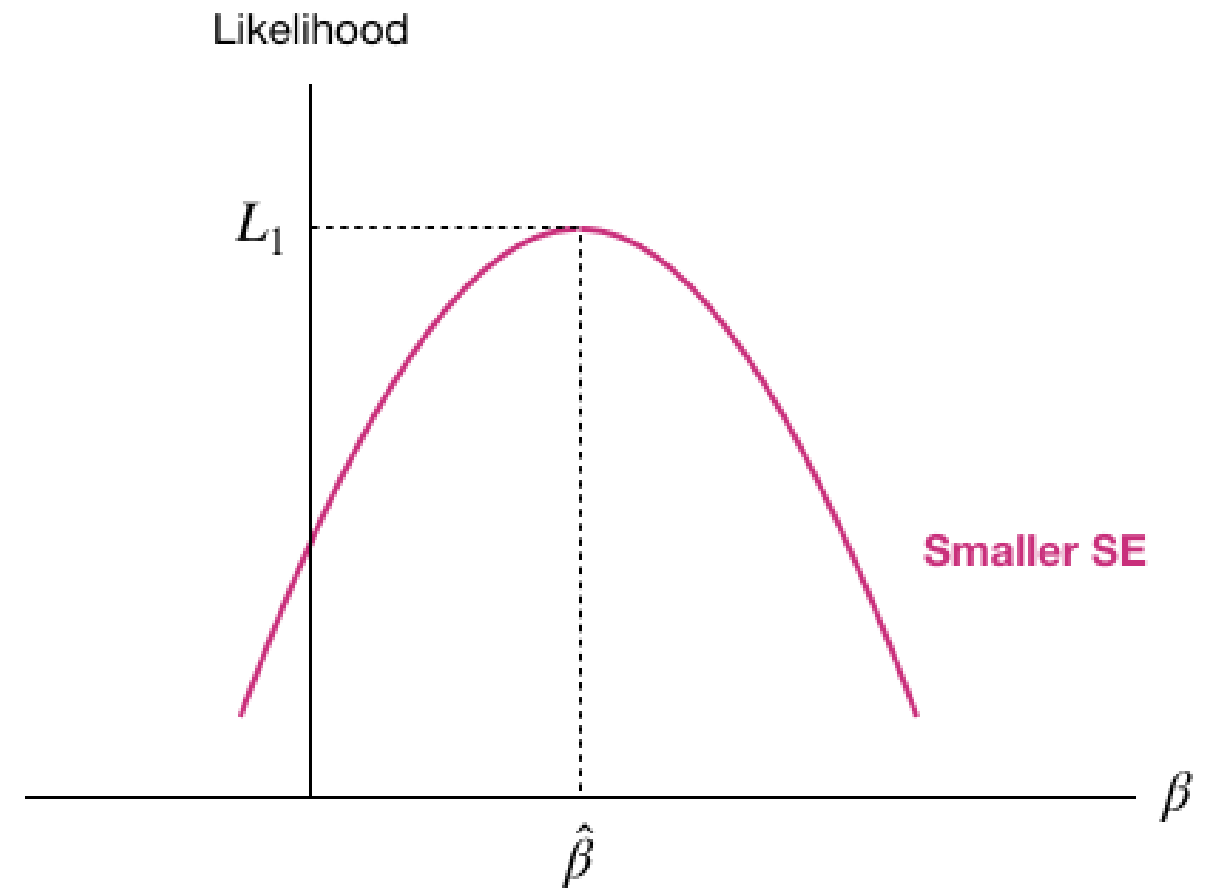


# Standard error (SE)

- Flatter peak
  - Location of maximum harder to define
  - Larger SE



- Sharper peak
  - Location of maximum more clearly defined
  - Smaller SE



# Computation of the standard error

```
# Extract variance-covariance matrix  
print(model_GLM.cov_params())
```

```
                Intercept    weight  
Intercept    0.774762 -0.325087  
weight      -0.325087  0.141903
```

```
# Compute standard error for weight  
std_error = np.sqrt(0.141903)
```

```
0.3767
```

## Variance-covariance matrix

	X	Y	Z
X	Variance X	Covariance X,Y	Covariance X,Z
Y	Covariance Y,X	Variance Y	Covariance Y,Z
Z	Covariance Z,X	Covariance Z,Y	Variance Z

# Significance testing

- z-statistic

$$z = \hat{\beta} / SE$$

- $z$  large  $\Rightarrow$  coefficient  $\neq 0 \Rightarrow$  variable significant
- Rule of thumb: cut-off value of 2

Example: horseshoe crab model

```
y ~ weight
```

$$z = 1.8151 / 0.377 = 4.819$$

# Confidence intervals for beta

- Uncertainty of the estimates
- 95% confidence intervals for  $\beta$

$$[\textit{lower}, \textit{upper}]$$

$$[\hat{\beta} - 1.96 \times SE, \hat{\beta} + 1.96 \times SE]$$

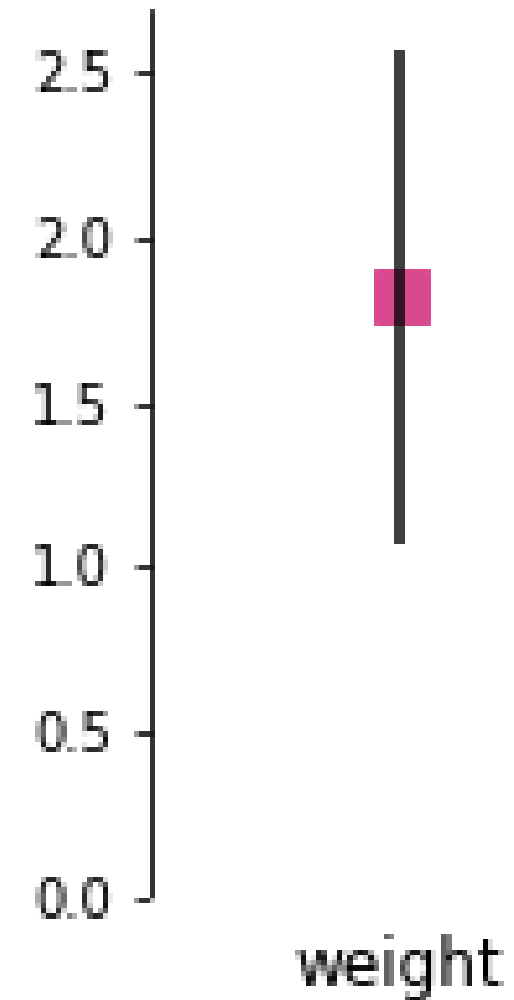
# Computing confidence intervals

Example: horseshoe crab model

	coef	std err
-----		
Intercept	-3.6947	0.880
weight	1.8151	0.377

$$[1.8151 - 1.96 \times 0.377, 1.8151 + 1.96 \times 0.377]$$

$$[1.07618, 2.55402]$$



# Extract confidence intervals

```
print(model_GLM.conf_int())
```

	0	1
Intercept	-5.419897	-1.969555
weight	1.076826	2.553463

# Extract confidence intervals

```
print(model_GLM.conf_int())
```

```
              lower      1  
Intercept -5.419897 -1.969555  
weight    1.076826  2.553463
```

# Extract confidence intervals

```
print(model_GLM.conf_int())
```

		0	upper
Intercept	-5.419897	-1.969555	
weight	1.076826	2.553463	



# Confidence intervals for odds

1. Extract confidence intervals for  $\beta$
2. Exponentiate endpoints

```
print(np.exp(model_GLM.conf_int()))
```

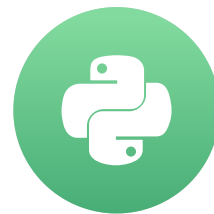
	0	1
Intercept	0.004428	0.139519
weight	2.935348	12.851533

# Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON

# Computing and describing predictions

GENERALIZED LINEAR MODELS IN PYTHON

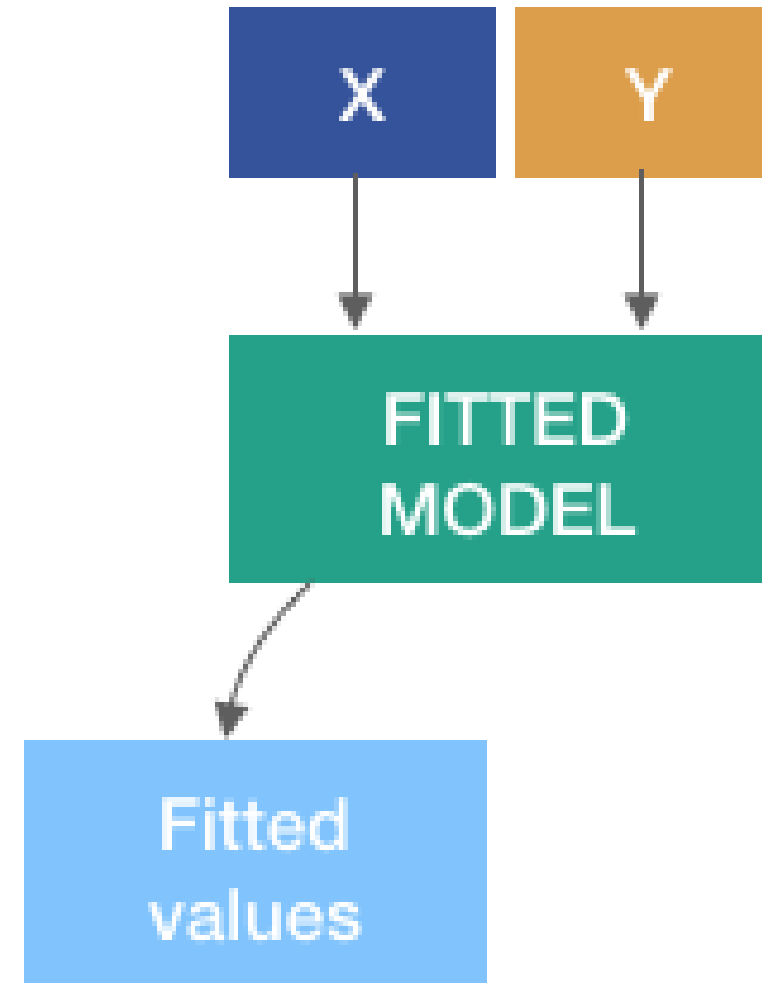


**Ita Cirovic Donev**  
Data Science Consultant

# Computing predictions

After obtaining model fit

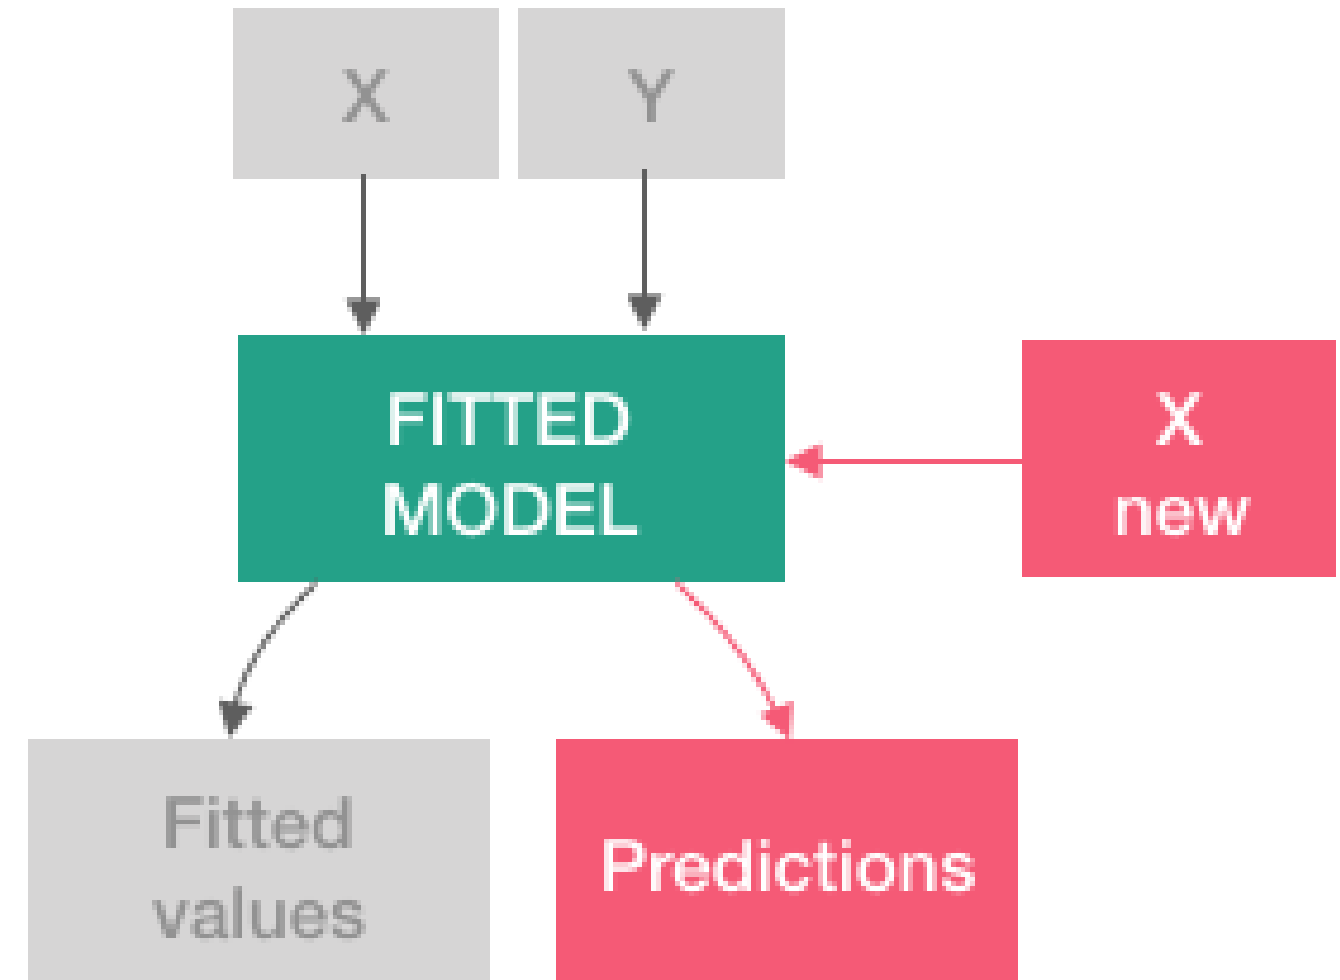
1. Fitted values for original  $x$  values



# Computing predictions

After obtaining model fit

1. **fitted values** for original  $x$  values
2. New values of  $x$  for **predicted values**



# Computing predictions

- Horseshoe crab model  $y \sim \text{weight}$

$$\mu = \frac{\exp(-3.6947 + 1.8151 \times \text{weight})}{1 + \exp(-3.6947 + 1.8151 \times \text{weight})}$$

- New measurement:  $\text{weight} = 2.85$

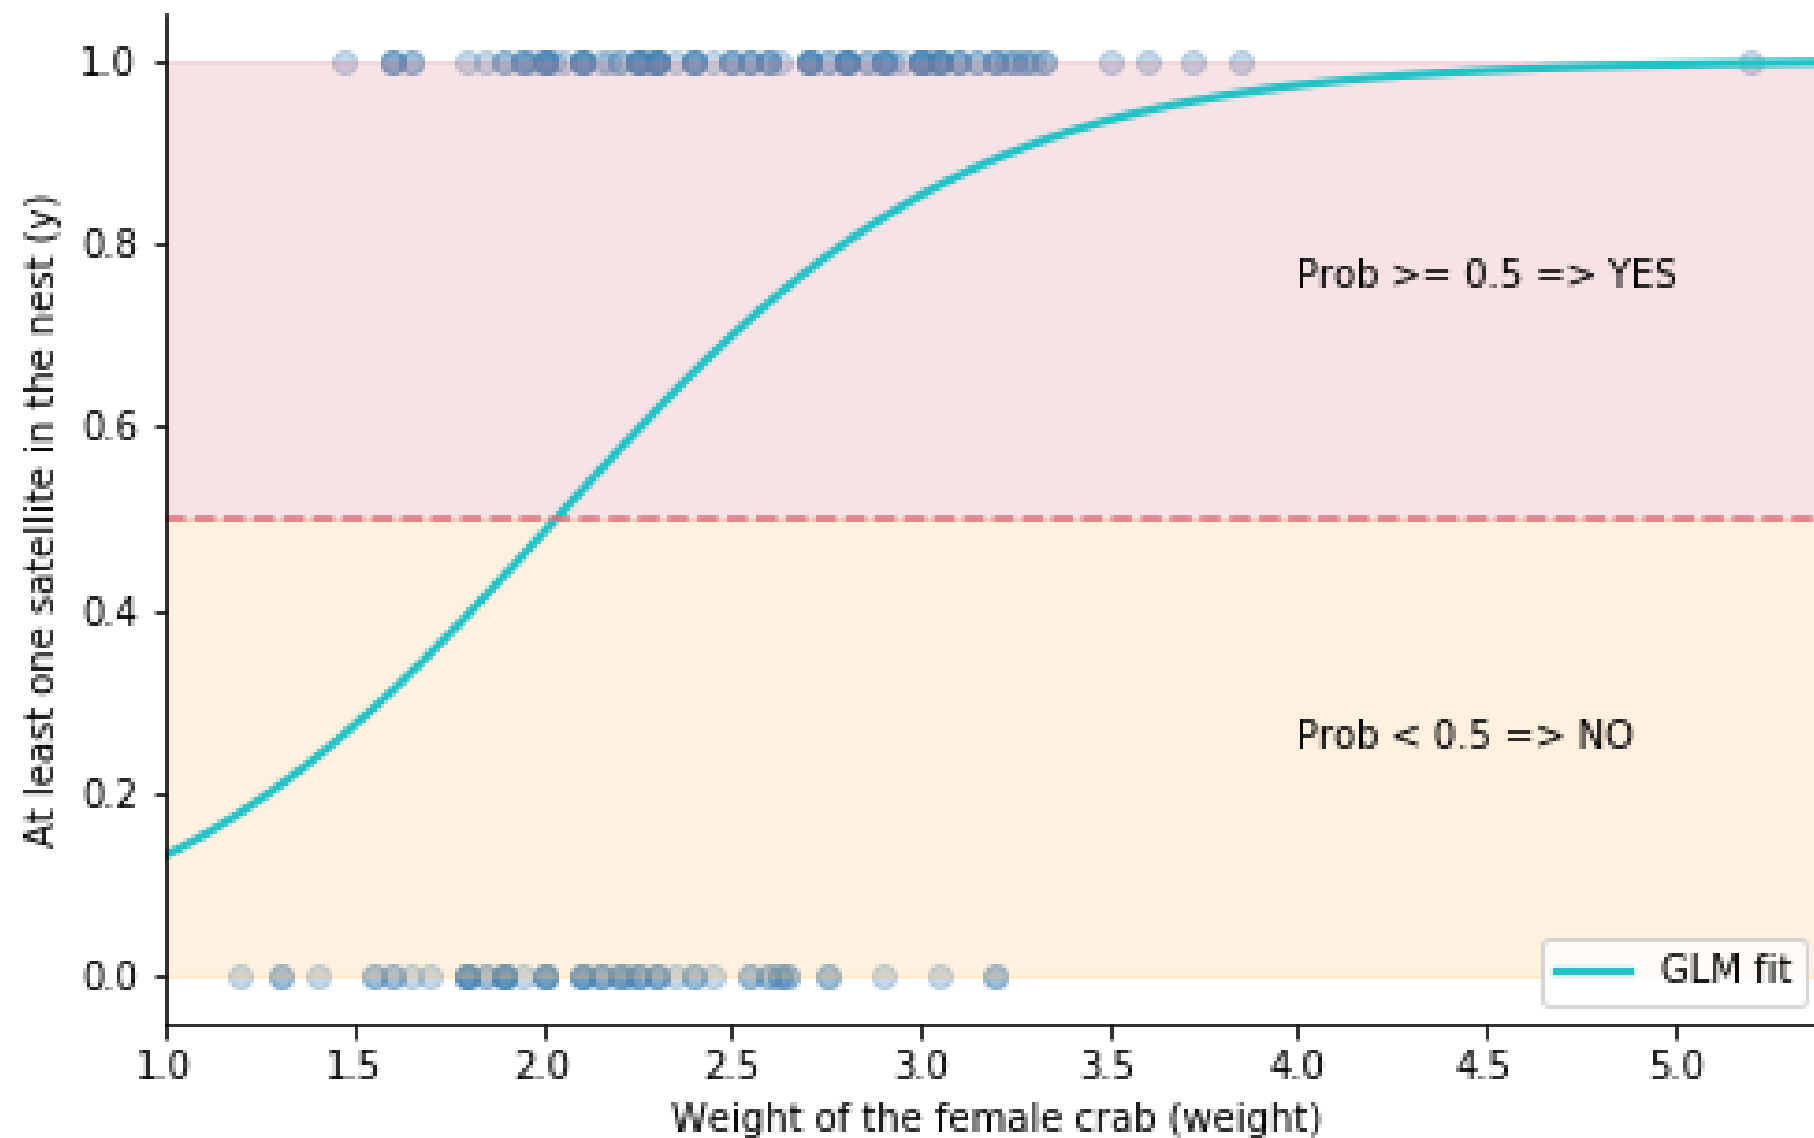
$$\mu = \frac{\exp(-3.6947 + 1.8151 \times 2.85)}{1 + \exp(-3.6947 + 1.8151 \times 2.85)} = 0.814$$

# Predictions in Python

- Compute model predictions for dataset `new_data`

```
# Compute model predictions  
model_GLM.predict(exog = new_data)
```

# From probabilities to classes





# Computing class predictions

```
# Extract fitted probabilities from model
crab['fitted'] = model.fittedvalues.values
```

```
# Define cut-off value
cut_off = 0.4
```

```
# Compute class predictions
crab['pred_class'] = np.where(crab['fitted'] > cut_off, 1, 0)
```

# Computing class predictions

```
# Count occurrences for each class  
crab['pred_class'].value_counts()
```

```
1    151  
0     22
```

Cut-off	$\hat{y} = 1$	$\hat{y} = 0$
$\mu = 0.4$	151	22
$\mu = 0.5$	126	47

# Confusion matrix

		Predicted	
		0	1
Actual	0		
	1		

# Confusion matrix - True Negatives

		Predicted	
		0	1
Actual	0	TN	
	1		

# Confusion matrix - True Positives

		Predicted	
		0	1
Actual	0	TN	
	1		TP

# Confusion matrix - False Positives

		Predicted	
		0	1
Actual	0	TN	FP
	1		TP

# Confusion matrix - False Negatives

		Predicted	
		0	1
Actual	0	TN	FP
	1	FN	TP

# Confusion matrix in Python

```
print(pd.crosstab(y_actual, y_predicted,  
                 rownames=['Actual'], colnames=['Predicted'],  
                 margins = True))
```

Predicted	0	1	All
Actual			
0	15	47	62
1	7	104	111
All	22	151	173



# Let's practice!

GENERALIZED LINEAR MODELS IN PYTHON